



Perf

Important Notice : Courseware - Legal

This courseware is both the product of the author and of freely available opensource and/or public domain materials. Wherever external material has been shown, it's source and ownership have been clearly attributed. We acknowledge all copyrights and trademarks of the respective owners.

The contents of the **courseware PDFs are considered proprietary** and thus cannot be copied or reproduced in any form whatsoever without the explicit written consent of the author.

Only the programs - **source code** and binaries (where applicable) - that form part of this courseware, and that are made available to the participant, are released under the terms of the [permissive MIT license](#).

Under the terms of the MIT License, you can certainly use the source code provided here; you must just attribute the original source (author of this courseware and/or other copyright/trademark holders).

VERY IMPORTANT :: Before using this source(s) in your project(s), you ***MUST*** check with your organization's legal staff that it is appropriate to do so.

The courseware PDFs are ***not*** under the MIT License, they are to be kept confidential, non-distributable without consent, for your private internal use only.

The duration, contents, content matter, programs, etc. contained in this courseware and companion participant VM are subject to change at any point in time without prior notice to individual participants.

Care has been taken in the preparation of this material, but there is no warranty, expressed or implied of any kind, and we can assume no responsibility for any errors or omissions. No liability is assumed for incidental or consequential damages in connection with or arising out of the use of the information or programs contained herein.

2000-2018 Kaiwan N Billimoria
kaiwanTECH, Bangalore, India.

kaiwanTECH Linux OS Corporate Training Programs
Please do check out our current offering of world-class, seriously-valuable, high on returns, technical Linux OS corporate training programs here: http://bit.ly/ktcorp

Perf is a multi-tool utility aimed at performance monitoring, cpu profiling and capture for later analysis. It is heavily maintained- mostly because it lives in the kernel source tree (under the *tools/* folder).

From `<kernel-src-tree>/tools/perf/design.txt` :

Performance Counters for Linux

Performance counters are special hardware registers available on most modern CPUs. These registers count the number of certain types of **hw events**: such as instructions executed, cache-misses suffered, or branches mis-predicted - without slowing down the kernel or applications. These registers can also trigger interrupts when a threshold number of events have passed - and can thus be used to profile the code that runs on that CPU.

The Linux Performance Counter subsystem provides an abstraction of these hardware capabilities. It provides per task and per CPU counters, counter groups, and it provides event capabilities on top of those. It provides "virtual" 64-bit counters, regardless of the width of the underlying hardware counters.

...

From the [Wikipedia entry on Perf](#):

perf (sometimes "Perf Events"[1] or perf tools, originally "Performance Counters for Linux", PCL[2]) - is a performance analyzing tool in Linux, available from kernel version 2.6.31 << July 2009 >>.[3] User-space controlling utility, called 'perf' has git-like interface with subcommands. It is capable of statistical profiling of entire system (both kernel and user code), single CPU or several threads.

It supports hardware performance counters, tracepoints, software performance counters (e.g. hrtimer), and dynamic probes (e.g. kprobes or uprobes).[4]

In 2012 IBM recognized perf (along with OProfile) as one of the two most commonly used performance counter profiling tools on Linux.[5]

...

Subcommands [\[edit\]](#)

perf is used with several subcommands:

- 'stat': measure total event count for single program or for system for some time
- 'top': [top](#)-like dynamic view of hottest functions
- 'record': measure and save sampling data for single program[6]
- 'report': analyze file generated by perf record; can generate flat, or graph profile.[6]
- 'annotate': annotate sources or assembly

- 'sched': tracing/measuring of scheduler actions and latencies[\[7\]](#)
- 'list': list available events

...

Perf is a profiler tool for Linux 2.6+ based systems **that abstracts away CPU hardware differences** in Linux performance measurements and presents a simple commandline interface. Perf is based on the *perf_events* interface exported by recent versions << **from ver 2.6.31** >> of the Linux kernel.

Installation

```
$ perf stat ps
```

```
WARNING: perf not found for kernel 3.16.0-38
```

You may need to install the following packages for this specific kernel:

```
linux-tools-3.16.0-38-generic
linux-cloud-tools-3.16.0-38-generic
```

You may also want to install one of the following packages to keep up to date:

```
linux-tools-generic
linux-cloud-tools-generic
```

```
$
```

Alternatively, install from Source:

On any recent (2.6.31 onwards) kernel source tree:

```
cd <root-of-kernel-src>
cd tools/perf
make
<< sudo cp ./perf /usr/bin >>
```

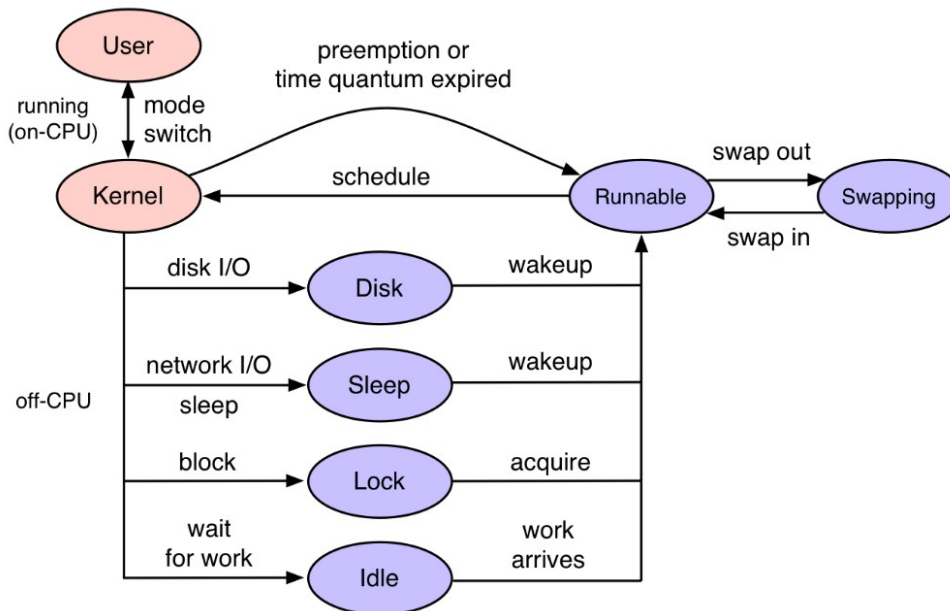
-OR- (easier)

```
sudo apt-get install linux-tools-$(uname -r)
```

Resources

- [Linux Perf by Brendan Gregg](#)
- [Linux Profiling - Perf](#)

- [A tutorial on Perf from the Perf wiki guide](#)
- [perf Examples](#)
- [Playing around with perf](#)
- [tools/perf/Documentation](#)
 - includes [tools/perf/Documentation/examples.txt](#)
- **IMP! Off-CPU Analysis**



- [Memory Access Patterns are Important](#)

Source

Events

perf_events instruments "events", which are a unified interface for different kernel instrumentation frameworks. The types of events are:

- **Hardware Events:** These instrument low-level processor activity based on CPU performance counters. For example, CPU cycles, instructions retired, memory stall cycles, level 2 cache misses, etc. Some will be listed as Hardware Cache Events.
- **Software Events:** These are low level events based on kernel counters. For example, CPU migrations, minor faults, major faults, etc.
- **Tracepoint Events:** These are kernel-level events based on the ftrace framework. These tracepoints are placed **in interesting and logical locations** of the kernel, so that higher-level behavior can be easily traced. For example, system calls, TCP events, file system I/O, disk I/O

O, etc. These are **grouped into libraries of tracepoints**; eg, "sock:" for socket events, "sched:" for CPU scheduler events.

Details about the events can be collected, including timestamps, the code path that led to it, and other specific details. The capabilities of perf_events are enormous, and you're likely to only ever use a fraction.

Apart from those categories of events, there are two other ways perf_events can instrument the system:

- **Profiling:** Snapshots can be **collected at an arbitrary frequency**, using `perf record -FHz`. This is commonly used **for CPU usage profiling**.
- **Dynamic Tracing:** Software can be **dynamically instrumented, creating events in any location**. For kernel software, this uses the **kprobes** framework. For user-level software, **uprobes**.

Tracepoints

List All available events that one can track:

Note!!

The capability and behaviour of perf varies with:

a) the hardware events supported by the processor

b) user privilege: running as root user shows all usable events; running as a regular user shows a subset of those

c) running perf on a virtual machine: likely you would have less events available than on a physical machine.

```
$ sudo /bin/bash
```

```
Password: <xxx>
```

```
# perf list
```

```
List of pre-defined events (to be used in -e):
```

cpu-cycles OR cycles	[Hardware event]
instructions	[Hardware event]
cache-references	[Hardware event]
cache-misses	[Hardware event]
branch-instructions OR branches	[Hardware event]
branch-misses	[Hardware event]
bus-cycles	[Hardware event]
stalled-cycles-frontend OR idle-cycles-frontend	[Hardware event]
stalled-cycles-backend OR idle-cycles-backend	[Hardware event]
ref-cycles	[Hardware event]
cpu-clock	[Software event]
task-clock	[Software event]
page-faults OR faults	[Software event]
context-switches OR cs	[Software event]
cpu-migrations OR migrations	[Software event]
minor-faults	[Software event]
major-faults	[Software event]
alignment-faults	[Software event]

Linux Debugging Techniques

Performance analysis with perf

```
emulation-faults
dummy
```

```
[Software event]
[Software event]
```

```
L1-dcache-loads
L1-dcache-load-misses
L1-dcache-stores
L1-dcache-store-misses
```

```
[Hardware cache event]
[Hardware cache event]
[Hardware cache event]
[Hardware cache event]
```

--snip--

```
exceptions:page_fault_kernel
exceptions:page_fault_user
```

```
[Tracepoint event]
[Tracepoint event]
```

<<

An example of the code within the kernel that 'counts' this perf event:

```
arch/mips/mm/fault.c
```

```
perf_sw_event(PERF_COUNT_SW_PAGE_FAULTS, 1, regs, address);
```

>>

--snip--

```
kvmmmu:fast_page_fault
kvmmmu:kvm_mmu_invalidate_zap_all_pages
kvmmmu:check_mmio_spte
kvm:kvm_entry
kvm:kvm_hypercall
kvm:kvm_hv_hypercall
kvm:kvm_pio
kvm:kvm_cpuid
```

```
[Tracepoint event]
[Tracepoint event]
[Tracepoint event]
[Tracepoint event]
[Tracepoint event]
[Tracepoint event]
[Tracepoint event]
[Tracepoint event]
```

--snip--

```
drm:drm_vblank_event_queued
drm:drm_vblank_event_delivered
skb:kfree_skb
skb:consume_skb
skb:skb_copy_datagram_iovec
net:net_dev_start_xmit
net:net_dev_xmit
net:net_dev_queue
net:netif_receive_skb
net:netif_rx
net:napi_gro_frags_entry
net:napi_gro_receive_entry
```

```
[Tracepoint event]
[Tracepoint event]
[Tracepoint event]
[Tracepoint event]
[Tracepoint event]
[Tracepoint event]
[Tracepoint event]
[Tracepoint event]
[Tracepoint event]
[Tracepoint event]
[Tracepoint event]
```

--snip--

```
block:block_split
block:block_bio_remap
block:block_rq_remap
jbd2:jbd2_checkpoint
jbd2:jbd2_start_commit
jbd2:jbd2_commit_locking
```

```
[Tracepoint event]
[Tracepoint event]
[Tracepoint event]
[Tracepoint event]
[Tracepoint event]
[Tracepoint event]
```

--snip--

```
fs:do_sys_open
fs:open_exec
```

```
[Tracepoint event]
[Tracepoint event]
```

```

migrate:mm_migrate_pages [Tracepoint event]
migrate:mm_numa_migrate_ratelimit [Tracepoint event]
compaction:mm_compaction_isolate_migratepages [Tracepoint event]
compaction:mm_compaction_isolate_freepages [Tracepoint event]
compaction:mm_compaction_migratepages [Tracepoint event]
compaction:mm_compaction_begin [Tracepoint event]
compaction:mm_compaction_end [Tracepoint event]
kmem:kmalloc [Tracepoint event]
kmem:kmem_cache_alloc [Tracepoint event]
kmem:kmalloc_node [Tracepoint event]
kmem:kmem_cache_alloc_node [Tracepoint event]
kmem:kfree [Tracepoint event]
kmem:kmem_cache_free [Tracepoint event]

--snip--

syscalls:sys_exit_unshare [Tracepoint event]
syscalls:sys_enter_mmap [Tracepoint event]
syscalls:sys_exit_mmap [Tracepoint event]
nmi:nmi_handler [Tracepoint event]
irq_vectors:local_timer_entry [Tracepoint event]
irq_vectors:local_timer_exit [Tracepoint event]

--snip--

xen:xen_cpu_write_gdt_entry [Tracepoint event]
xen:xen_cpu_set_ldt [Tracepoint event]
#

# perf list 'kmem:*'
kmem:kmalloc [Tracepoint event]
kmem:kmem_cache_alloc [Tracepoint event]
kmem:kmalloc_node [Tracepoint event]
kmem:kmem_cache_alloc_node [Tracepoint event]
kmem:kfree [Tracepoint event]
kmem:kmem_cache_free [Tracepoint event]
kmem:mm_page_free [Tracepoint event]
kmem:mm_page_free_batched [Tracepoint event]
kmem:mm_page_alloc [Tracepoint event]
kmem:mm_page_alloc_zone_locked [Tracepoint event]
kmem:mm_page_pcpu_drain [Tracepoint event]
kmem:mm_page_alloc_extfrag [Tracepoint event]
#

```

Summarizing the tracepoint “library names” and numbers of tracepoints, on my system:

```

# perf list | awk -F: '/Tracepoint event/ { lib[$1]++ } END {
    for (l in lib) { printf "%-16s %d\n", l, lib[l] } }' | sort | column
block          19      iwlwifi_io      7      rcu          1
btrfs          40      iwlwifi_msg    5      regmap       15
cfg80211       137     iwlwifi_ucode  2      regulator    7
compaction     5       jbd2           16     rpm          4
context_tracking 2      kmem           12     sched        23
drm            3       kvm            42     scsi          5
exceptions     2      kvmmmu         14     signal        2

```

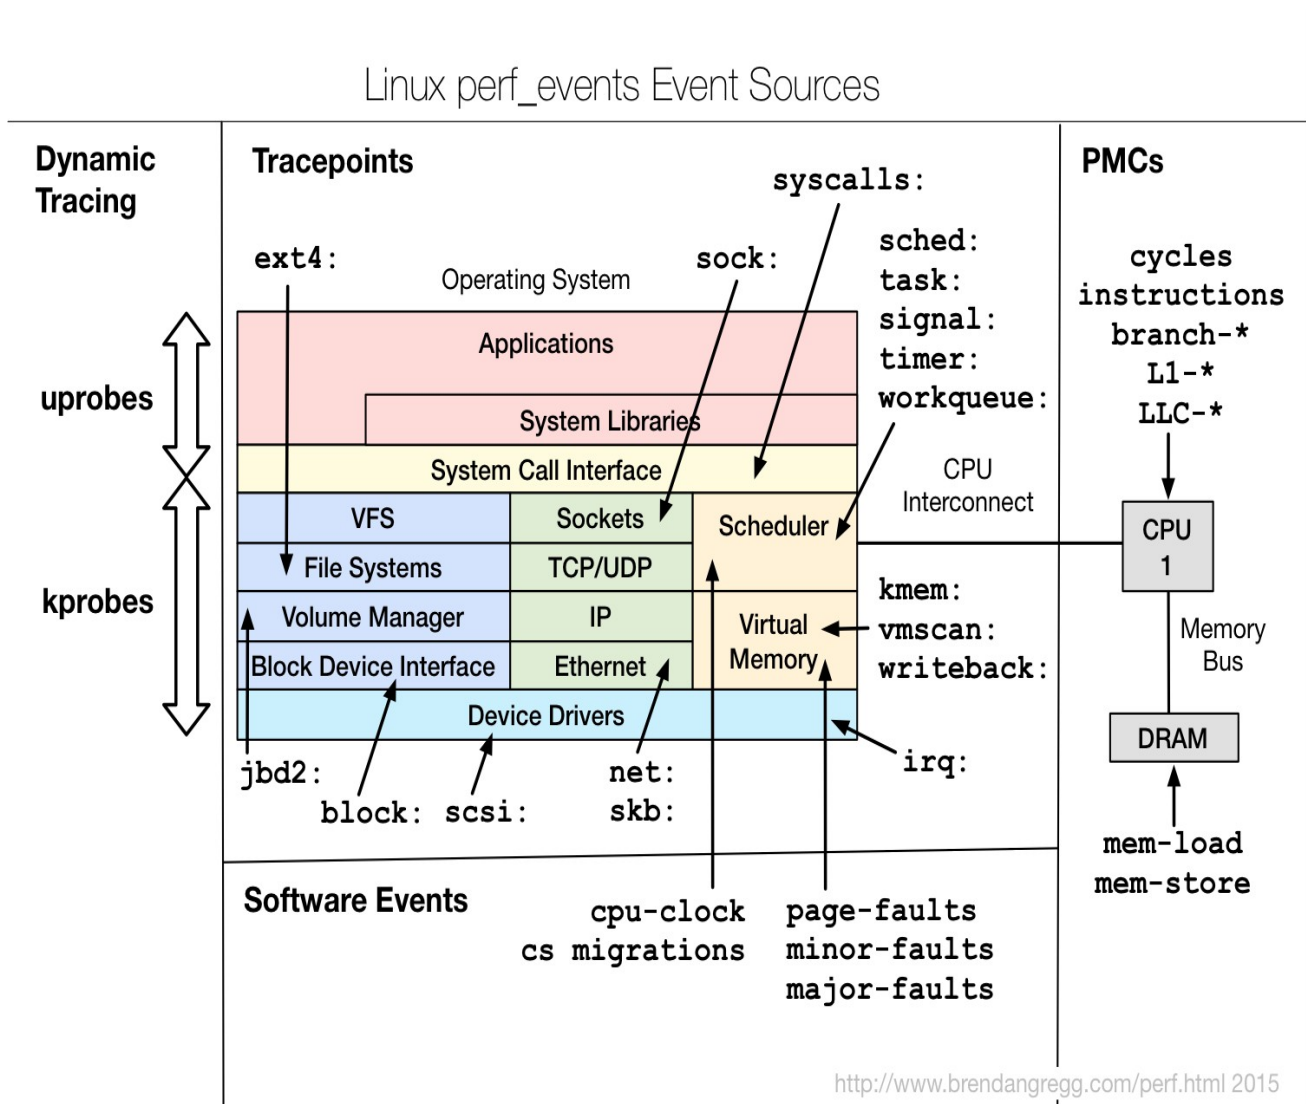

ext4	91	mac80211	101	skb	3
filelock	6	mac80211_msg	3	sock	2
filemap	2	mce	1	spi	7
fs	2	migrate	2	swiotlb	1
ftrace	1	module	5	syscalls	588
gpio	2	napi	1	task	2
hda	7	net	10	timer	13
hda_intel	2	nmi	1	udp	1
i2c	8	oom	1	vmscan	15
i915	29	pagemap	2	vsyscall	1
iommu	7	power	21	workqueue	4
irq	5	printk	1	writeback	25
irq_vectors	20	random	15	xen	35
iwlwifi	5	ras	1	xfs	348
iwlwifi_data	2	raw_syscalls	2	xhci-hcd	9
#					

These include:

- **block:** block device I/O
- **ext3, ext4:** file system operations
- **kmem:** kernel memory allocation events
- **random:** kernel random number generator events
- **sched:** CPU scheduler events
- **syscalls:** system call enter and exits
- **task:** task events

It's worth checking the list of tracepoints after every kernel upgrade, to see if any are new. The value of adding them [has been debated](#) from time to time, with it wondered if anyone will use them (I will!). There is a balance to aim for: I'd include the smallest number of probes that sufficiently covers common needs, and anything unusual or uncommon can be left to dynamic tracing.

[P.T.O.]

All events diagram:

[PMC = Performance Measurement Counter]

Also:

<< [Source](#) >> ...

Preflight check

First thing to do when you start working with Perf is to launch **perf test**. This will check your system and kernel features and report if something isn't available. Usually, you need to make as much as possible «OK"s.

Beware though that perf will behave differently when launched under «root» and ordinary user. It's smart enough to let you do some things without root privileges. There is a control file at `/proc/sys/kernel/perf_event_paranoid` that you can tweak in order to change access to perf events:

```
$ perf stat -a
```

```
Error:
```

```
You may not have permission to collect system-wide stats.
```

Consider tweaking `/proc/sys/kernel/perf_event_paranoid`:

- 1 - Not paranoid at all
- 0 - Disallow raw tracepoint access for unpriv
- 1 - Disallow cpu events for unpriv
- 2 - Disallow kernel profiling for unpriv

After you played with *perf test*, you can see what hardware events are available to you with `perf list`. Again, the list will differ depending on current user id. Also, the amount of events will depend on your hardware: x86_64 CPUs have much more hardware events than some low-end ARM processors.

...

Simple Example

```
# perf stat -a -e 'kmem:kmalloc' sleep 1    << -a => system wide; -e =>
tracepoint event >>
```

Performance counter stats for 'system wide':

```
5,347          kmem:kmalloc
```

1.001556007 seconds time elapsed

#

```
# perf stat -e kmem:* sleep 1    << only for the 'sleep' process >>
```

Performance counter stats for 'sleep 1':

```
6          kmem:kmalloc
53         kmem:kmem_cache_alloc
0          kmem:kmalloc_node
0          kmem:kmem_cache_alloc_node
15         kmem:kfree
55         kmem:kmem_cache_free
43         kmem:mm_page_free
40         kmem:mm_page_free_batched
38         kmem:mm_page_alloc
6          kmem:mm_page_alloc_zone_locked
6          kmem:mm_page_pcpu_drain
0          kmem:mm_page_alloc_extfrag
```

1.001691044 seconds time elapsed

#

Tip: Leave out the '-a' option to report the number of `kmalloc`'s that occurred for precisely the process that was run ('sleep 1' in the above example).

Top-like system-wide Profiling

To get a quick view of overall system state, do:

perf top

```

root@Seawolf-VirtAppliance: ~
File Edit Tabs Help
root@Seaw... x root@Seaw... x
Samples: 1K of event 'cpu-clock', Event count (approx.): 123743603
Overhead Shared Object Symbol
68.72% vboxvideo_drv_117.so [.] 0x00000000000006e81
5.66% [vdso] [.] 0x0000000000000949
4.54% [kernel] [k] native_read_tsc
4.39% [kernel] [k] _raw_spin_unlock_irqrestore
3.62% [kernel] [k] finish_task_switch
1.46% [kernel] [k] tick_nohz_idle_enter
0.48% libgobject-2.0.so.0.4600.1 [.] g_type_check_instance_is_a
0.41% libc-2.21.so [.] __memcpy_sse2_unaligned
0.34% [kernel] [k] retint_careful
0.32% libglib-2.0.so.0.4600.1 [.] g_hash_table_lookup
0.31% [kernel] [k] __fget
0.31% libglib-2.0.so.0.4600.1 [.] g_mutex_lock
0.31% libgtk-3.so.0.1600.7 [.] 0x0000000000179825
0.26% [kernel] [k] do_setitimer
0.24% libc-2.21.so [.] __strcmp_sse2_unaligned
0.20% Xorg [.] 0x00000000001bf2c3
0.20% libcairo.so.2.11400.2 [.] 0x0000000000155d7
0.20% libglib-2.0.so.0.4600.1 [.] 0x0000000000065d80
0.20% libgobject-2.0.so.0.4600.1 [.] 0x00000000001a9a0
0.20% libpango-1.0.so.0.3600.8 [.] pango_default_break
0.20% libpangoft2-1.0.so.0.3600.8 [.] pango_fc_font_lock_face
0.20% libvte.so.9.2800.2 [.] 0x0000000000295a8
0.20% perf [.] 0x000000000006702c
0.20% [kernel] [k] __do_softirq
0.20% [kernel] [k] entry_SYSCALL_64_after_swapgs
0.19% perf [.] 0x0000000000085fa4
For a higher level overview, try: perf top --sort comm,dso

```

The above screenshot was taken on a Linux appliance running on VirtualBox. As can be clearly seen, the majority of cpu overhead is that of the VirtualBox video driver (almost 70%!).

Now keep the highlight bar (grey colour), on the top line and press right-arrow key: the following shows up:

So we now see that this runs in the process context of Xorg pid 703, i.e., the X server process of course.

```

Zoom into Xorg(703) thread
Zoom into vboxvideo_drv_117.so DS0
Browse map details
Run scripts for samples of thread [Xorg]
Run scripts for all samples
Run scripts for all samples
Exit

```

Another example: on an Ubuntu VM, some of the Intel e1000 NIC driver's 'methods' seem to be banging on CPU:

Samples: 974 of event 'cpu-clock', Event count (approx.): 11622035		
Overhead	Shared Object	Symbol
19.88%	[kernel]	[k] e1000_xmit_frame
16.37%	[kernel]	[k] e1000_alloc_rx_buffers
14.70%	[kernel]	[k] e1000_watchdog
12.98%	[kernel]	[k] _raw_spin_unlock_irqrestore
9.59%	[kernel]	[k] __softirqentry_text_start
6.72%	[kernel]	[k] tick_nohz_idle_enter
6.30%	[kernel]	[k] e1000_clean
2.15%	libslang.so.2.3.1	[.] SLtt_smart_puts

Samples: 164 of event 'cpu-clock', Event count (approx.): 9634037		
Overhead	Shared Object	Symbol
22.26%	[kernel]	[k] e1000_xmit_frame
18.39%	[kernel]	[k] e1000_watchdog
15.60%	[kernel]	[k] _raw_spin_unlock_irqrestore
14.86%	[kernel]	[k] e1000_alloc_rx_buffers
7.96%	[kernel]	[k] tick_nohz_idle_enter
6.50%	[kernel]	[k] __softirqentry_text_start
3.07%	[kernel]	[k] e1000_clean
1.99%	libslang.so.2.3.1	[.] 0x0000000000009affb
1.99%	perf	[.] 0x000000000001e7864
1.16%	perf	[.] 0x000000000001ab96c

Quick Summary by process (context):

perf top --sort comm,dso

Samples: 3K of event 'cpu-clock', Event count (approx.): 550103079		
Overhead	Command	Shared Object
25.78%	:1709	[kernel]
24.07%	Xorg	vboxvideo_drv_117.so
14.33%	:1703	[kernel]
9.15%	:0	[kernel]
8.47%	:1709	libbz2.so.1.0.4
5.24%	:1703	libapt-pkg.so.4.16.0
4.00%	:1709	libapt-pkg.so.4.16.0
1.32%	:1702	[kernel]
1.10%	:1656	[kernel]
1.05%	Xorg	[kernel]
0.52%	Xorg	Xorg
0.50%	:1699	[kernel]

TIP- to see PID:name, shared object and symbolic name simultaneously:

```
$ sudo perf top -r 90 --sort pid,comm,dso,symbol << -r 90 => collect
data with SCHED_FIFO RT scheduling class and priority 90 [1-99] >>
```

Samples: 1M of event 'cycles', Event count (approx.): 259022836877

Overhead	Pid:Command	Command	Shared Object	Symbol
1.63%	17213:perf	perf	perf	[.] 0x0000000000023a0a6
0.64%	17213:perf	perf	perf	[.] 0x000000000019fa07
0.18%	17213:perf	perf	perf	[.] 0x00000000001e8304
0.17%	17210:perf	perf	perf	[.] 0x00000000001ab0ed
0.15%	17210:perf	perf	perf	[.] 0x00000000001e896f
0.15%	2095:Xorg	Xorg	Xorg	[.] 0x000000000013d16f
0.14%	7891:python2	python2	python2.7	[.] PyEval_EvalFrameEx
0.14%	2095:Xorg	Xorg	[kernel]	[k] fw_domains_get
0.13%	3772:compiz	compiz	libcompiz_core.so.0.9.13.1	[.] CompWindow::destroyed
0.12%	17210:perf	perf	perf	[.] 0x000000000019fa07
0.11%	3772:compiz	compiz	libopengl.so	[.] PrivateGLScreen::paintOutp
0.11%	17213:perf	perf	perf	[.] 0x00000000001dfe5
0.10%	2095:Xorg	Xorg	Xorg	[.] 0x000000000013906b
0.10%	3772:compiz	compiz	[kernel]	[k] fw_domains_get
0.09%	17213:perf	perf	perf	[.] 0x000000000019f1c5
0.09%	7507:firefox	firefox	libxcb.so.1.1.0	[.] 0x000000000000cc55
0.09%	17213:perf	perf	perf	[.] 0x0000000000249be7
0.09%	17210:perf	perf	libc-2.24.so	[.] __strcmp_sse2_unaligned
0.08%	17210:perf	perf	perf	[.] 0x00000000001dfe8f
0.07%	7507:firefox	firefox	libpthread-2.24.so	[.] pthread_mutex_lock
0.07%	2095:Xorg	Xorg	libfb.so	[.] fbGetWindowPixmap
0.07%	17213:perf	perf	perf	[.] 0x00000000001e8267
0.07%	2095:Xorg	Xorg	Xorg	[.] 0x0000000000ef5c6
0.07%	7507:firefox	firefox	libxul.so	[.] 0x00000000009d43cd
0.07%	17213:perf	perf	libc-2.24.so	[.] __strcmp_sse2_unaligned
0.07%	3772:compiz	compiz	libc-2.24.so	[.] __int_malloc
0.07%	17213:perf	perf	perf	[.] 0x00000000001e9596
0.06%	17213:perf	perf	perf	[.] 0x00000000001dfea8

no symbols found in /home/kaiwan/MentorGraphics/Sourcery CodeBench Lite for ARM GNU Linux/bin/arm-none-linux-gnueabi-ld, mayb

For convenience

```
$ alias ptop
alias ptop='sudo perf top'
$ alias ptopv
alias ptopv='sudo perf top -r 80 -f 99 --sort pid,comm,dso,symbol \
--demangle-kernel -v --call-graph dwarf,fractal'
$
```

Q. How does one get a report?

A. 2 steps:

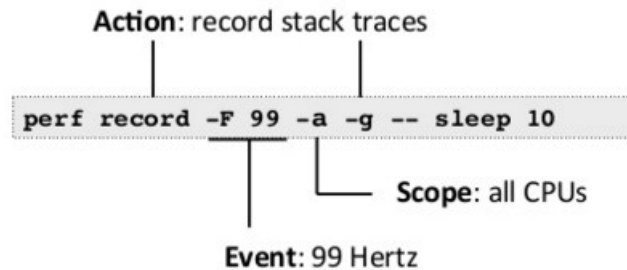
1. Use the

```
perf record -e <events> [ -F <sampling_frequency> -a ] [command]
```

syntax to capture data to a binary file (called *perf.data* by default, use *-output <fname>* to save to an alternate pathname).

perf Command Format

- perf instruments using `stat` or `record`. This has three main parts: action, event, scope.
- e.g., profiling on-CPU stack traces:



Note: `sleep 10` is a dummy command to set the duration

<< Note: substitute `'-g'` with `'--call-graph dwarf'` >>

perf Scope

- System-wide: all CPUs (`-a`)
- Target PID (`-p PID`)
- Target command (...)
- Specific CPUs (`-c ...`)
- User-level only (`<event>:u`)
- Kernel-level only (`<event>:k`)
- A custom filter to match variables (`--filter ...`)

2. Use the

perf report [`--stdio`]

syntax to display & analyze the report.

<<

SIDEBAR

Perf for Android?

This page from Linaro - [Using Perf on Android](#) - covers it quite well.

>>

Eg. <See the “Profiling” and “Static Tracing” section [here](#) for many examples >

Sample CPU stack traces for the entire system, at 99 Hertz, for 10 seconds:

perf record -F 99 -a --call-graph dwarf sleep 10

-F 99 => frequency of 99 (i.e. sample @ a rate of 99 times / second)

-a => scope: all cpu's

--call-graph dwarf (older: -g) => Enables call-graph (stack chain/backtrace) recording. (More details follow).

perf report --stdio

Failed to open /tmp/perf-3465.map, continuing without symbols

To display the perf.data header info, please use --header/--header-only options.

#

Samples: 2K of event 'cycles'

Event count (approx.): 49250631497

#

Children Self Command Shared Object

.....
.....

#

5.83% 0.00% firefox [unknown] [...] 0000000000000000

|
--- (nil)

|
|--16.38%-- 0x0

|
|
|--92.00%-- 0x7f7a198351c4
| 0xf88948fffa8f17e9

|
|--8.00%-- 0x0

|
|--14.00%-- 0x10200
| 0x7573736920746f6e

|

--snip--

4.45% 0.00% evince-thumbnai [kernel.kallsyms] [k] system_call_fastpath

|
--- system_call_fastpath

|
|--40.15%-- open64

|
|
|--96.40%-- _dl_map_object

|
|--3.60%-- open_path
| _dl_map_object

|
|--8.30%-- mprotect

|
|
|--65.98%-- _dl_map_object

|
|--34.02%-- dl_main
| _dl_sysdep_start

|

--snip--

#

Example 1 : What's causing system boot to take so long?

On a custom product (based on Ubuntu 15.10), boot was taking (much) longer than expected. How does one investigate this?

I setup a cron job (as root) to let perf sample and record system-wide:

```
# crontab -l
...
@reboot /usr/bin/perf record -F 512 -a --call-graph dwarf --output
</path/to/perf-out> -s -T --realtime=99
#

[ -F 512 : record @ freq of 512 HZ
  -a      : all cpu's, full system
  --call-graph dwarf : record stack using DWARF
  --output : specify perf data output file (default: perf.data)
  -s      : per-thread counts
  -T      : sample timestamps (see with perf script)
  --realtime=99 : run perf with scheduling policy SCHED_FIFO max
priority (99)
                    (root access required)
]
```

After booting:

```
# perf report --input </path/to/perf-out> -f
```

Samples: 2K of event 'cycles', Event count (approx.): 22247721974					
Children	Self	Command	Shared Object	Symbol	
+ 16.88%	0.06%	apt-check	python3.4	[.]	PyEval_EvalCodeEx
+ 16.88%	0.29%	apt-check	python3.4	[.]	PyEval_EvalFrameEx
+ 15.81%	0.00%	apt-check	libc-2.21.so	[.]	__libc_start_main
+ 15.81%	0.00%	apt-check	python3.4	[.]	main
+ 15.81%	0.00%	apt-check	python3.4	[.]	Py_Main
+ 15.79%	0.00%	apt-check	python3.4	[.]	_start
+ 15.49%	0.00%	apt-check	python3.4	[.]	PyEval_EvalCode
+ 15.47%	0.00%	apt-check	python3.4	[.]	PyRun_FileExFlags
+ 15.47%	0.00%	apt-check	python3.4	[.]	0x00000000001f7452
+ 15.45%	0.00%	apt-check	python3.4	[.]	PyRun_SimpleFileExFlags
+ 10.76%	0.00%	freshclam	libclamav.so.6.1.26	[.]	0x000000000042d52
+ 10.76%	0.00%	freshclam	libclamav.so.6.1.26	[.]	0x0000000000041c20
+ 10.76%	0.00%	freshclam	libz.so.1.2.8	[.]	gzread
+ 10.46%	0.00%	freshclam	libz.so.1.2.8	[.]	0x000000000010a7a
+ 10.46%	0.06%	freshclam	libz.so.1.2.8	[.]	inflate
+ 10.38%	0.00%	apt-check	python3.4	[.]	PyObject_Call
+ 9.02%	0.00%	apt-check	python3.4	[.]	0x0000000000159c55
+ 8.57%	0.10%	swapper	[kernel.kallsyms]	[k]	cpu_startup_entry
+ 7.88%	0.00%	apt-check	apt_pkg.cpython-34m-x86_64-linux-gnu.so	[.]	0x00000000002064f
+ 7.88%	0.00%	apt-check	libapt-pkg.so.4.16.0	[.]	_ZN12pkgCacheFile4OpenEP100pProgress
+ 7.67%	0.00%	apt-check	libapt-pkg.so.4.16.0	[.]	_ZN12pkgCacheFile13BuildDepCacheEP100pProgress
+ 7.67%	0.30%	apt-check	libapt-pkg.so.4.16.0	[.]	_ZN11pkgDepCache4InitEP100pProgress
+ 6.87%	0.01%	swapper	[kernel.kallsyms]	[k]	call_cpuidle
+ 6.85%	0.00%	swapper	[kernel.kallsyms]	[k]	cpuidle_enter
+ 5.55%	0.75%	apt-check	libapt-pkg.so.4.16.0	[.]	_ZN11pkgDepCache6UpdateEP100pProgress
+ 4.76%	0.00%	apt-check	python3.4	[.]	0x0000000000cf143
+ 4.73%	0.00%	apt-check	python3.4	[.]	0x0000000000159cac
+ 4.63%	0.00%	swapper	[kernel.kallsyms]	[k]	x86_64_start_kernel

The screenshot above shows that 'apt-check' and 'freshclam' are consuming large amounts of cpu. (apt-check, indeed, [has been known to consume cpu at boot](#). Disabling it helps!

Regarding freshclam/clamav – it's an opensource antivirus solution – there are two choices: either uninstall clamav or reconfigure freshclam to have the daemon work in the background).

After reconfiguring freshclam and rebooting the system, the perf report shows:

```
# perf report -i perf_all.data --sort comm,dso
```

```
+ 28.49%    2.90% apt-check      python3.4
+ 27.39%    3.95% apt-check      libc-2.21.so
+ 26.52%   22.00% apt-check      libapt-pkg.so.4.16.0
+ 24.89%    0.45% apt-check      apt_pkg.cpython-34m-x86_64-linux-
gnu.so
+  8.12%    8.12% swapper        [kernel.kallsyms]
+  5.17%    0.99% Xorg          Xorg
+  4.85%    0.31% Xorg          libc-2.21.so
+  3.99%    0.41% zenity         libgtk-3.so.0.1600.7
+  3.81%    1.03% lxpanel        libc-2.21.so
+  3.47%    0.01% lxpanel        lxpanel
+  3.24%    0.46% pcmanfm        libc-2.21.so
...
...
```

Disabling apt-check at boot then eliminates it's overhead.

[BTW there are interesting tools like [bootchart](#) too, a visualization tool to show boot-time activity (it renders a Gantt-like chart displaying which processes are taking cpu)].

<<

*Hey, guess what, there's an easier (and very detailed) way to see 'who took so damn long' to startup – if you're using **systemd**:*

```
$ systemd-analyze <tab-tab>
```

```
blame          critical-chain  dot          dump          plot
set-log-level  syscall-filter  time        verify
$
```

Eg.

CLI: `systemd-analyze blame`

GUI: `systemd-analyze plot` -writes an SVG!

Eg.

```
$ systemd-analyze blame
```

```
32.050s apt-daily-upgrade.service
30.276s plymouth-quit-wait.service
27.698s apt-daily.service
14.548s snapd.service
13.839s dev-sda1.device
11.632s ModemManager.service
10.680s udisks2.service
 9.950s dev-loop10.device
 9.904s dev-loop11.device
 9.857s dev-loop12.device
 9.839s dev-loop6.device
...
```

```
6.252s apparmor.service
6.209s apport.service
5.481s NetworkManager-wait-online.service
5.336s polkit.service
4.648s avahi-daemon.service
4.543s systemd-udev.service
4.378s wpa_supplicant.service

...

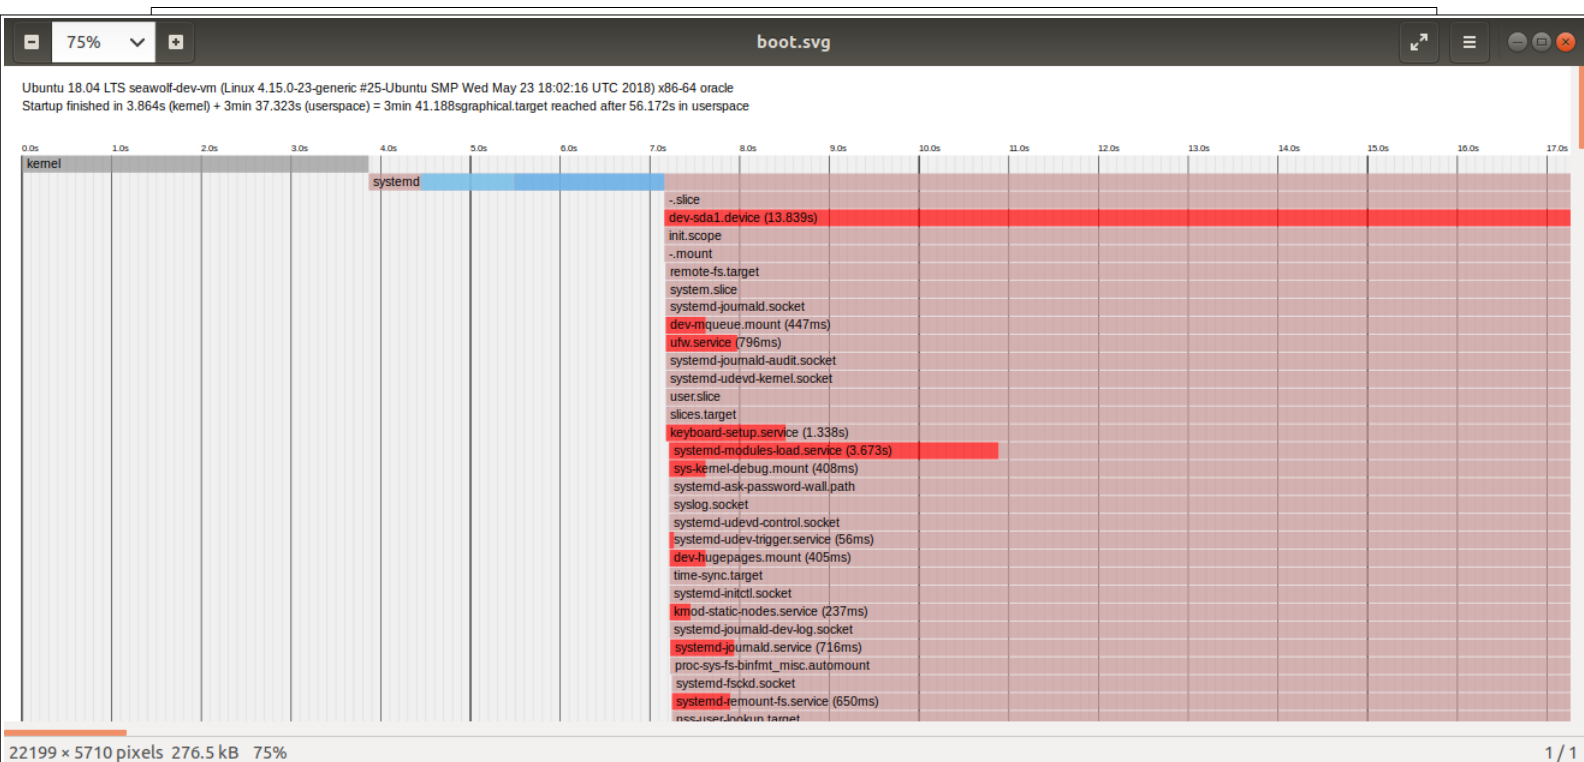
28ms dev-loop17.device
19ms snapd.socket
13ms systemd-update-utmp-runlevel.service
6ms sys-kernel-config.mount
4ms sys-fs-fuse-connections.mount

$
```

GUI: `systemd-analyze plot` -writes an SVG!

```
$ systemd-analyze plot > boot-time.svg
$
```

View the SVG within a web browser (or image viewer app):



Very interestingly, once [snap](#) was purged from Ubuntu 18.04 LTS, the startup time reduced dramatically:

	Kernel	Userspace
With snap	3.864s	3min 31.323s
Without snap	2.906s	26.191s

>>

Example 2 : firefox hanging

On my Ubuntu 15.10 box, firefox seemed to be hung up. Tools like top or htop could not reveal much. (Recall, capturing the “Who? and How?” are easy, but the “What? and Why?” are not so easy!).

So, lets use perf:

```
# pgrep firefox << get firefox's PID >>
26189
# perf record --call-graph dwarf -p 26189 << record stats including stack >>
Lowering default frequency rate to 3250.
```

```
Please consider tweaking /proc/sys/kernel/perf_event_max_sample_rate.
```

```
...
```

```
^C[ perf record: Woken up 15 times to write data ]
```

```
[ perf record: Captured and wrote 8.671 MB perf.data (557 samples) ]
```

```
#
```

```
#
```

```
# ls -lh perf.data
```

```
-rw----- 1 root root 8.7M Dec 29 13:11 perf.data
```

```
# perf report --stdio
```

```
[P.T.O.]
```

```
# To display the perf.data header info, please use --header/--header-only options.
#
# Total Lost Samples: 0
#
# Samples: 557 of event 'cycles'
# Event count (approx.): 207872625
#
# Children      Self  Command      Shared Object      Symbol
# .....
#
# 63.43%    0.00%  DOM Worker    libc-2.21.so        [.] __clone
#           |
#           ---__clone
#
# 63.43%    0.00%  DOM Worker    libpthread-2.21.so  [.] start_thread
#           |
#           ---start_thread
#           ___clone
#
# 63.43%    0.00%  DOM Worker    libnspr4.so         [.] 0x00000000000296c8
#           |
#           ---0x7f2b0abe86c8
#           start_thread
#           ___clone
#
# 63.43%    0.00%  DOM Worker    libxul.so            [.] 0x000000000009e042
#           |
#           ---0x7f2b05428442
#           0x7f2b0abe86c8
#           start_thread
#           ___clone
#
# 63.43%    0.00%  DOM Worker    libxul.so            [.] 0x00000000000bbc962
#           |
#           ---0x7f2b05604962
#           0x7f2b05428442
```

Clearly, firefox seems to be stuck in spawning threads (notice the multiple “___clone”s in the stack traces).

Furthermore, a quick check with ps shows that “DOM Worker” is a child thread of the firefox process:

```
$ ps -LA|grep "^26189" |wc -l          << ps: -L options shows threads >>
```

```
81
```

```
$ ps -LA|grep "^26189"
```

```
26189 26189 ?          01:50:24 firefox
26189 26292 ?          00:00:00 gmain
26189 26293 ?          00:00:11 gdbus
26189 26294 ?          00:00:00 dconf worker
26189 26295 ?          00:00:07 Gecko_IOThread
26189 26296 ?          00:00:00 Link Monitor
26189 26297 ?          00:06:59 Socket Thread
26189 26298 ?          00:00:00 firefox
26189 26299 ?          00:01:11 JS Helper
```

```

26189 26300 ?      00:01:12 JS Helper
26189 26301 ?      00:01:14 JS Helper
26189 26302 ?      00:01:15 JS Helper
26189 26303 ?      00:01:13 JS Helper
26189 26304 ?      00:01:13 JS Helper
26189 26305 ?      00:01:12 JS Helper
26189 26306 ?      00:01:15 JS Helper
26189 26307 ?      00:00:01 JS Watchdog
26189 26308 ?      00:00:00 Hang Monitor
26189 26312 ?      00:00:15 Cache2 I/O
26189 26313 ?      00:05:11 Timer
26189 26314 ?      00:00:00 firefox
26189 26315 ?      00:00:00 Cert Verify
26189 26317 ?      00:00:00 GMPThread
26189 26319 ?      00:03:16 Compositor
26189 26320 ?      00:00:23 ImageBridgeChil
26189 26321 ?      00:00:14 ImgDecoder #1
26189 26322 ?      00:00:13 ImgDecoder #2
26189 26323 ?      00:00:13 ImgDecoder #3
26189 26324 ?      00:00:00 ImageIO
26189 26325 ?      00:05:03 SoftwareVsyncTh
26189 26326 ?      00:00:01 HTML5 Parser
26189 26327 ?      00:00:00 IPDL Background
26189 26331 ?      00:00:14 DOM Worker
26189 26335 ?      00:00:21 DOM Worker
26189 26341 ?      00:00:00 DOM Worker
26189 26342 ?      00:00:01 mozStorage #1
26189 26343 ?      00:00:00 firefox
26189 26344 ?      00:00:00 Cache I/O
...
26189 14708 ?      00:00:00 firefox
$

```

Found them! But which thread of the several (below) “DOM Worker” threads above is the culprit? Lets look at individual threads:

```

$ ps -LA|grep "^26189.*DOM Worker"
26189 26331 ?      00:00:14 DOM Worker
26189 26335 ?      00:00:21 DOM Worker
26189 26341 ?      00:00:00 DOM Worker
26189 26375 ?      00:00:13 DOM Worker
26189 1573 ?       00:00:30 DOM Worker
26189 12012 ?      00:00:37 DOM Worker
$

```

To see the report by thread, use the `--tid` switch:

```

$ perf report --stdio --tid=26331
# To display the perf.data header info, please use --header/--header-
only options.
#
#
# Total Lost Samples: 0

```

```
#
# Samples: 557 of event 'cycles'
# Event count (approx.): 207872625
#
# Children      Self    Command      Shared Object      Symbol
# .....      .....      .....      .....      .....
#
#      1.23%      0.00%  DOM Worker  libxul.so          [...]
0x00000000029de5b0
|
---0x7f2b074265b0
    0x7f2b072284e0
    0x7f2b074852b5
    0x7f2b0748617e
    0x7f2b07486522
    0x7f2b074883d0
    0x7f2b0647f23f
    0x7f2b0647f29e
    0x7f2b0648786b
    0x7f2b0648d167
    0x7f2b0648ef7e
    0x7f2b06467a3e
    0x7f2b05425268
    0x7f2b05441667
    0x7f2b056146bd
    0x7f2b05604962
    0x7f2b05428442
    0x7f2b0abe86c8
    start_thread
    __clone
...
...
$
```

Does not seem to be the thread causes CPU overhead (only 1.23% overhead). So we check out the other “DOM Worker” threads until we find the one eating CPU:

...

```
$ perf report --stdio --tid=1573 -v
  symsrc__init: cannot get elf header.
Looking at the vmlinux_path (7 entries long)
Using /proc/kcore for kernel object code
Using /proc/kallsyms for symbols
unwind: pthread_cond_timedwait@@GLIBC_2.3.2:ip = 0x7f2b0bf75149 (0xd149)
unwind: reg 16, val 7f2b0bf75149
unwind: reg 7, val 7f2ae82f2b00
unwind: find_proc_info dso /lib/x86_64-linux-gnu/libpthread-2.21.so
unwind: no map for 7f2b0c1812b0
unwind: access_mem 0x7f2b0c1812b0 not inside range 0x7f2ae82f2b00-0x7f2ae82f4000
unwind: pthread_cond_timedwait@@GLIBC_2.3.2:ip = 0x7f2b0bf75149 (0xd149)
...
```

```

...
unwind: __clone:ip = 0x7f2b0b1ffeed (0x106eed)
# To display the perf.data header info, please use --header/--header-
only options.
#
#
# Total Lost Samples: 0
#
# Samples: 557 of event 'cycles'
# Event count (approx.): 207872625
#
# Children      Self   Command      Shared Object
Symbol
# .....
.. .....
#
   63.43%      0.00%  DOM Worker  /lib/x86_64-linux-gnu/libc-2.21.so
0x106eed      u [.] __clone
      |
      ---__clone

   63.43%      0.00%  DOM Worker  /lib/x86_64-linux-gnu/libpthread-
2.21.so 0x76aa      1 [.] start_thread
      |
      ---start_thread
      __clone

   63.43%      0.00%  DOM Worker  /usr/lib/firefox/libnspr4.so
0x296c8      1 [.] 0x000000000000296c8
      |
      ---0x7f2b0abe86c8
      start_thread
      __clone

   63.43%      0.00%  DOM Worker  /usr/lib/firefox/libxul.so
0x9e0442      1 [.] 0x000000000009e0442
      |
      ---0x7f2b05428442
      0x7f2b0abe86c8
      start_thread
      __clone

   63.43%      0.00%  DOM Worker  /usr/lib/firefox/libxul.so
0xbbc962      1 [.] 0x0000000000bbc962
      |
      ---0x7f2b05604962
      0x7f2b05428442
      0x7f2b0abe86c8
      start_thread
      __clone

   63.43%      0.00%  DOM Worker  /usr/lib/firefox/libxul.so
0xbcc6bd      1 [.] 0x0000000000bcc6bd
      |

```



```

    ---0x7f2b056146bd
    0x7f2b05604962
    0x7f2b05428442
    0x7f2b0abe86c8
    start_thread
    __clone
...
$

```

Example 3 : Qemu/KVM: Running an x86_64 (Debian) guest with and without KVM hypervisor acceleration

Environment:

Host: Ubuntu 17.04, 4.10.0-35-generic kernel.

Host CPU:

```
$ lscpu | egrep -i "name|vt|ept"
```

```
Model name:      Intel(R) Core(TM) i7-6500U CPU @ 2.50GHz
```

```
Virtualization:  VT-x
```

```
Flags:           fpu vme [...] flexpriority ept vpid [...]
```

```
$
```

qemu-system-x86_64 full virtualization:

```
$ qemu-system-x86_64 --version
```

```
QEMU emulator version 2.8.0 (Debian 1:2.8+dfsg-3ubuntu2.5)
```

```
Copyright (c) 2003-2016 Fabrice Bellard and the QEMU Project developers
```

```
$
```

A helper script

```
[...]
```

```
QEMU_ARGS_COMMON="-kernel ${KERNEL_IMG} -drive file=${ROOTFS},if=virtio,format=raw -append root=/dev/vda"
```

```
[...]
```

```
# Perf stuff
```

```
PERF_REC_ARGS="-F 256 -s -T --realtime=99 --call-graph dwarf"
```

```
PERF_RAWDATA_WITHKVM=perf_raw_withkvm.data
```

```
PERF_RAWDATA_NOKVM=perf_raw_nokvm.data
```

```
[...]
```

3.1 WITH KVM

<< First, we make sure the *kvm.ko* and *kvm-[intel|amd].ko* kernel modules are loaded up >>

```
...
```

```

qemu-system-x86_64 --enable-kvm ${QEMU_ARGS_COMMON} \
    -chardev stdio,id=stdio,mux=on,signal=off \
    -device virtio-serial-pci -device virtio-net-pci \

```

```

-device virtconsole,chardev=stdio \
-net nic,model=virtio \
-mon chardev=stdio \
-fsdev local,id=fs1,path=${
{SHARED_HOST_FOLDER}},security_model=none \
-device virtio-9p-pci,fsdev=fs1,mount_tag=host-code & << run in
background >>

```

```

qpid=$(pidof -s ${QEMUPRCS})
[ -z "${qpid}" ] && {
echo "Failed to get PID of ${QEMUPRCS}, abandoning perf run.."
} || {
HZ=512
sudo perf record -p ${qpid} ${PERF_REC_ARGS} --output $
{PERF_RAWDATA_WITHKVM}
}

```

<< After the run, perf record saves the raw data file (3 MB) >>

```
$ sudo perf report -i perf_raw_withkvm.data
```

```

Samples: 333 of event 'cycles', Event count (approx.): 1995020328

```

Children	Self	Command	Shared Object	Symbol
+ 61.10%	0.00%	qemu-system-x86	libc-2.24.so	[.] __clone
+ 61.10%	0.00%	qemu-system-x86	libpthread-2.24.so	[.] start_thread
+ 59.89%	0.06%	qemu-system-x86	[kernel.kallsyms]	[k] entry_SYSCALL_64_fastpath
- 58.88%	0.00%	qemu-system-x86	qemu-system-x86_64	[.] kvm_cpu_exec
-	-	kvm_cpu_exec		
-	-	41.37% kvm_vcpu_ioctl		
-	-	GI_ioctl		
-	-	39.20% entry_SYSCALL_64_fastpath		
-	-	37.77% sys_ioctl		
-	-	37.69% do_vfs_ioctl		
-	-	kvm_vcpu_ioctl		
-	-	34.95% kvm_arch_vcpu_ioctl_run		
-	-	28.86% vcpu_enter_guest		
-	-	20.72% vmx_handle_exit		
+	+	9.98% handle_ept_misconfig		
+	+	3.40% handle_cpuid		
+	+	2.97% handle_ept_violation		
+	+	1.65% handle_apic_access		
+	+	1.36% handle_io		
+	+	2.79% vmx_save_host_state		
+	+	1.81% __srcu_read_unlock		
+	+	1.16% __srcu_read_lock		
+	+	1.69% kvm_vcpu_block		
+	+	0.86% __srcu_read_lock		
+	+	0.85% vmx_handle_exit		
+	+	0.77% complete_emulated_mmio		
+	+	0.66% _cond_resched		
+	+	2.70% vcpu_put		
+	+	1.37% syscall_return_slowpath		
+	+	0.61% vmx_vcpu_run		
+	+	17.07% address_space_write		
+ 58.88%	0.00%	qemu-system-x86	qemu-system-x86_64	[.] 0xfffffa9cad3195918
+ 42.52%	0.00%	qemu-system-x86	libc-2.24.so	[.] __GI_ioctl
+ 41.37%	0.00%	qemu-system-x86	qemu-system-x86_64	[.] kvm_vcpu_ioctl

Majority of time spent by qemu-system-x86_64 is in kvm_cpu_exec → kvm_vcpu_ioctl → ...

3.2 WITHOUT KVM

<< First, we make sure the `kvm.ko` and `kvm-[intel|amd].ko` kernel modules are *removed* >>

```
[...]
qemu-system-x86_64 ${QEMU_ARGS_COMMON} & << run in background >>

qpid=$(pidof -s ${QEMU_PRCS})
[ -z "${qpid}" ] && {
    echo "Failed to get PID of ${QEMU_PRCS}, abandoning perf run.."
} || {
    HZ=99
    sudo perf record -p ${qpid} ${PERF_REC_ARGS} --output
      ${PERF_RAWDATA_NOKVM}
}

[ perf record: Captured and wrote 102.977 MB perf_raw_nokvm.data (12763
samples) ]
$ ls -lh perf_raw_nokvm.data
-rw----- 1 root root 103M Sep 30 12:15 perf_raw_nokvm.data
$
```

We can literally “feel” that it’s much much slower!

```
$ sudo perf report -i perf_raw_nokvm.data
```

Samples: 12K of event 'cycles', Event count (approx.): 227457053992

	Children	Self	Command	Shared Object	Symbol
+	81.72%	0.00%	qemu-system-x86	libc-2.24.so	[.] __clone
+	81.72%	0.00%	qemu-system-x86	libpthread-2.24.so	[.] start_thread
+	81.14%	5.59%	qemu-system-x86	qemu-system-x86_64	[.] cpu_exec
+	81.14%	0.00%	qemu-system-x86	qemu-system-x86_64	[.] 0xfffffaa79cd0ff491
-	61.10%	1.30%	qemu-system-x86	qemu-system-x86_64	[.] tb_gen_code
-	59.80%		tb_gen_code		
-	48.95%		tcg_gen_code		
	11.41%		tcg_optimize		
+	1.09%		0x2d018d		
	1.01%		0x2cf01c		
	0.85%		0x2cf9f8		
	0.56%		0x2cedc9		
	0.52%		0x2cf779		
	0.51%		0x2cf5a8		
+	8.56%		gen_intermediate_code		
+	1.46%		qht_insert		
+	1.30%		__clone		
-	48.95%	20.31%	qemu-system-x86	qemu-system-x86_64	[.] tcg_gen_code
-	28.64%		tcg_gen_code		
	11.41%		tcg_optimize		
+	1.09%		0x2d018d		
	1.01%		0x2cf01c		
	0.85%		0x2cf9f8		
	0.56%		0x2cedc9		
	0.52%		0x2cf779		
	0.51%		0x2cf5a8		
+	20.31%		__clone		
+	13.41%	0.37%	qemu-system-x86	qemu-system-x86_64	[.] _init
+	11.41%	10.12%	qemu-system-x86	qemu-system-x86_64	[.] tcg_optimize
+	8.56%	0.25%	qemu-system-x86	qemu-system-x86_64	[.] gen_intermediate_code
+	7.80%	0.38%	qemu-system-x86	qemu-system-x86_64	[.] get_page_addr_code
+	5.72%	0.02%	qemu-system-x86	qemu-system-x86_64	[.] tlb_fill
+	5.70%	0.34%	qemu-system-x86	qemu-system-x86_64	[.] x86_cpu_handle_mmu_fault
+	5.43%	4.53%	qemu-system-x86	qemu-system-x86_64	[.] qht_lookup
+	3.85%	3.57%	qemu-system-x86	qemu-system-x86_64	[.] iotlb_to_region

This time, the majority of time spent by `qemu-system-x86_64` is in the **TCG** (Tiny Code Generator – a JIT compiler) code paths: `tcg_gen_code` → ... , implying that it's Qemu's dynamic binary translation engine. Makes sense, as we now can't run guest code directly on the host cpu and instead have to rely on - much slower – binary translation techniques.

Example 4 : calling `fork()` / `exit()` in a loop 60,000 times

Children	Self	Command	Shared Object
- 96.67%	96.67%	fork test	[kernel.kallsyms]
- 81.33%		__GI__fork (inlined)	
- 69.81%		entry_SYSCALL_64_after_hwframe	
-		do_syscall_64	
- 69.50%		_do_fork	
- 65.97%		copy_process.part.35	
- 17.39%		copy_page_range	
+ 5.63%		__pte_alloc	
+ 3.11%		__pmd_alloc	
+ 2.17%		__pud_alloc	
+ 6.98%		anon_vma_fork	
+ 6.34%		__vmalloc_node_range	
+ 5.94%		perf_event_init_task	
+ 3.38%		kmem_cache_alloc	
+ 2.66%		security_vm_enough_memory_mm	
1.98%		dup_userfaultfd	
+ 1.44%		copy_fs_struct	
+ 1.33%		arch_dup_task_struct	
+ 1.32%		mm_init	
+ 1.18%		sched_fork	

Partial screenshot above: we can see that the majority of the work is in `copy_process()` and within that in the function `copy_page_range()`; within that, its the *page table setup* code that takes time.

<< *Related Tutorial: [“PERF tutorial: Finding execution hot spots”](#)* >>

To get the most out **perf**, you'll want **symbols and stack traces**. These may work by default in your Linux distribution, or they may require the addition of packages, or recompilation of the kernel with additional config options.

Symbols

`perf_events`, like other debug tools, **needs symbol information (symbols)**. These are used to translate memory addresses into function and variable names, so that they can be read by us humans. Without symbols, you'll see hexadecimal numbers representing the memory addresses profiled.

...

<<

On a Fedora 26 box

`perf` crashed, with error message “*** Error in `perf': free(): invalid pointer: 0x0000559531a30d20 ***”. So, we fire up GDB to examine the corefile :

```
$ gdb -q -c
corefile:host=localhost.fc26:gPID=24912:gTID=24912:ruid=1000:sig=6:exe=\!
usr\!bin\!perf.24912 /usr/bin/perf
Reading symbols from /usr/bin/perf... (no debugging symbols found) ...done.
[New LWP 24912]
```

```
warning: Unexpected size of section `.reg-xstate/24912' in core file.
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib64/libthread_db.so.1".
Core was generated by `perf report --stdio -T --input=perf.data'.
Program terminated with signal SIGABRT, Aborted.
```

```
warning: Unexpected size of section `.reg-xstate/24912' in core file.
#0  0x00007ff95772a69b in raise () from /lib64/libc.so.6
Missing separate debuginfos, use: dnf debuginfo-install perf-4.13.5-200.fc26.x86_64
```

<< Interesting! Fedora asks us to install the appropriate dbgsym ver of perf! >>

```
(gdb) bt
#0  0x00007ff95772a69b in raise () from /lib64/libc.so.6
#1  0x00007ff95772c4a0 in abort () from /lib64/libc.so.6
#2  0x00007ff9577708e1 in __libc_message () from /lib64/libc.so.6
#3  0x00007ff95777b789 in _int_free () from /lib64/libc.so.6
#4  0x00007ff9577810ee in free () from /lib64/libc.so.6
#5  0x000055a6069265b5 in perf_read_values_destroy ()
#6  0x000055a60687bde1 in cmd_report ()
#7  0x000055a6068dc8b1 in ?? ()
#8  0x000055a6068dcbae in ?? ()
#9  0x000055a606864130 in main ()
```

```
(gdb)
```

```
>>
```

Another thing that can happen when not running perf as root:

```
$ perf record -F 200 -a sleep 3
```

```
WARNING: Kernel address maps (/proc/{kallsyms,modules}) are restricted,
check /proc/sys/kernel/kptr_restrict.
```

Samples in kernel functions may not be resolved if a suitable vmlinux file is not found in the buildid cache or in the vmlinux path.

Samples in kernel modules won't be resolved at all.

If some relocation was applied (e.g. kexec) symbols may be misresolved even with a suitable vmlinux or kallsyms file.

Cannot read kernel map

Error:

You may not have permission to collect system-wide stats.
Consider tweaking /proc/sys/kernel/perf_event_paranoid:

- 1 - Not paranoid at all
- 0 - Disallow raw tracepoint access for unpriv
- 1 - Disallow cpu events for unpriv
- 2 - Disallow kernel profiling for unpriv

```
$
```

Stack Traces

Always compile with frame pointers. **Omitting frame pointers is an evil compiler optimization that breaks debuggers, and sadly, is often the default.** Without them, you may see incomplete stacks

from `perf_events`, like seen in the earlier `sshd` symbols example. There are two ways to fix this: either using dwarf data to unwind the stack, or returning the frame pointers.

Note!

For compiled code (for development time at least), always use the “`-fno-omit-frame-pointer`” gcc switch !!

Dwarf ¹

Since about the 3.9 kernel, `perf_events` has supported a **workaround for missing frame pointers** in user-level stacks: `libunwind`, which uses dwarf. This can be enabled using “`-g dwarf`”.

Eg. # Sample CPU stack traces for the PID, using dwarf to unwind stacks, at 99 Hertz, for 10 seconds:
`perf record -F 99 -p PID -g dwarf sleep 10`

<<

Common Error:

```
# perf record -F 99 -g dwarf -p 2869
Workload failed: No such file or directory
#
```

See [this post](#).

“... That error is telling you have a **bad command line**. I remember running into this error in the past and basically if `perf` fails to parse the command line it will give this nondescript “Workload failed: No such file or directory” error in some cases. It's since been fixed in new versions of `perf`.

In your first example your passing in: ``-g dwarf`` which is wrong. If you read the man page that comes up when you run `perf help` report you'll see that the ``-g`` flag doesn't take any parameters. Instead, you have to specify that after ``--call-graph FOO``, like in your second example, which takes the callgraph method as a parameter. Also, ``--call-graph FOO`` already implies `-g`. ”

Ok, bottom-line: don't use the arg as “`-g dwarf`”, instead use “`--call-graph dwarf`”:

```
# perf record -F 99 --call-graph dwarf -p 2869
<< ...wait a few seconds... >>
^C[ perf record: Woken up 154 times to write data ]
[ perf record: Captured and wrote 38.802 MB perf.data (~1695291
samples) ]
#
>>

...
```

If user-level stacks look incomplete, you can try `perf record` with “~~`-g dwarf`~~” “`--call-graph dwarf`” as a different technique to unwind them. See the [Stacks](#) section.

The output from `perf report` can be many pages long, which can become cumbersome to read. Try

¹ What exactly is the DWARF format and how can one see it? See this [article here](#).

generating [Flame Graphs](#) from the same data.

...

One-Liners: Counting Events

```
# CPU counter statistics for the specified command:
perf stat command

# Detailed CPU counter statistics (includes extras) for the specified command:
perf stat -d command

# CPU counter statistics for the specified PID, until Ctrl-C:
perf stat -p PID

# CPU counter statistics for the entire system, for 5 seconds:
perf stat -a sleep 5

# Various CPU last level cache statistics for the specified command:
perf stat -e LLC-loads,LLC-load-misses,LLC-stores,LLC-prefetches command

# Count system calls for the specified PID, until Ctrl-C:
perf stat -e 'syscalls:sys_enter_*' -p PID

# Count scheduler events for the specified PID, for 10 seconds:
perf stat -e 'sched:*' -p PID sleep 10

# Count block device I/O events for the entire system, for 10 seconds:
perf stat -e 'block:*' -a sleep 10

# Show system calls by process, refreshing every 2 seconds:
perf top -e raw_syscalls:sys_enter -ns comm
```



```
# Sample on-CPU functions for the specified command, at 99 Hertz:
perf record -F 99 command

# Sample on-CPU functions for the specified PID, at 99 Hertz, until Ctrl-C:
perf record -F 99 -p PID

# Sample CPU stack traces for the specified PID, at 99 Hertz, for 10 seconds:
perf record -F 99 -p PID -g -- sleep 10

# Sample CPU stack traces for the entire system, at 99 Hertz, for 10 seconds:
perf record -F 99 -ag -- sleep 10

# Sample CPU stack traces, once every 10,000 Level 1 data cache misses, for 5 s:
perf record -e L1-dcache-load-misses -c 10000 -ag -- sleep 5

# Sample CPU stack traces, once every 100 last level cache misses, for 5 seconds:
perf record -e LLC-load-misses -c 100 -ag -- sleep 5

# Sample on-CPU kernel instructions, for 5 seconds:
perf record -e cycles:k -a -- sleep 5

# Sample on-CPU user instructions, for 5 seconds:
perf record -e cycles:u -a -- sleep 5
```

Interesting:

Look carefully at the (5th and) 6th example above:

```
perf record -e LLC-load-misses -c 100 -a --call-graph dwarf --
sleep 5
```

LLC: Last Level Cache.

Meaning: for all CPUs system-wide, record stats for 5 sec. Every time we get 100 cache load misses, sample the cpu stack trace. This helps in profiling apps that have cache misses. Fine tune for your app by using the -p (PID) switch.

<<

Monitoring software for cache misses can be useful. What perf events are available for this?

Eg. on a laptop with an Intel Core-i7 processor:

```
# perf list|grep -i miss
cache-misses [Hardware event]
branch-misses [Hardware event]
L1-dcache-load-misses [Hardware cache event]
L1-dcache-store-misses [Hardware cache event]
L1-dcache-prefetch-misses [Hardware cache event]
L1-icache-load-misses [Hardware cache event]
dTLB-load-misses [Hardware cache event]
dTLB-store-misses [Hardware cache event]
iTLB-load-misses [Hardware cache event]
branch-load-misses [Hardware cache event]
branch-misses OR cpu/branch-misses/ [Kernel PMU event]
cache-misses OR cpu/cache-misses/ [Kernel PMU event]
```

#

Example:# **perf stat -e cache-misses vim /etc/passwd**

Performance counter stats for 'vim /etc/passwd':

2,07,480 cache-misses

3.042336654 seconds time elapsed

#

>>

Reporting

```
# Show perf.data in an ncurses browser (TUI) if possible:
perf report

# Show perf.data with a column for sample count:
perf report -n

# Show perf.data as a text report, with data coalesced and percentages:
perf report --stdio

# List all raw events from perf.data:
perf script

# List all raw events from perf.data, with customized fields:
perf script -f comm,tid,pid,time,cpu,event,ip,sym,dso

# Dump raw contents from perf.data as hex (for debugging):
perf script -D

# Disassemble and annotate instructions with percentages (needs some debuginfo):
perf annotate --stdio
```

Before getting into the how of a FlameGraph:

<< from [here](#) and [here](#) : slides of Brendan Gregg's talk at Xscale, Feb 2015 >>

CPU Workload Characterization

- Who
- Why
- What
- How

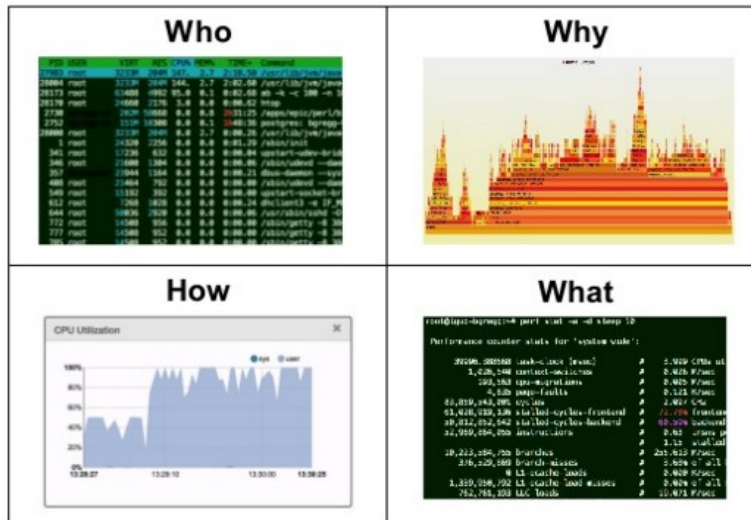
- For CPUs:

1. **Who:** which PIDs, programs, users
2. **Why:** code paths, context
3. **What:** CPU instructions, cycles
4. **How:** changing over time

- Can you currently answer them? How?



CPU Tools



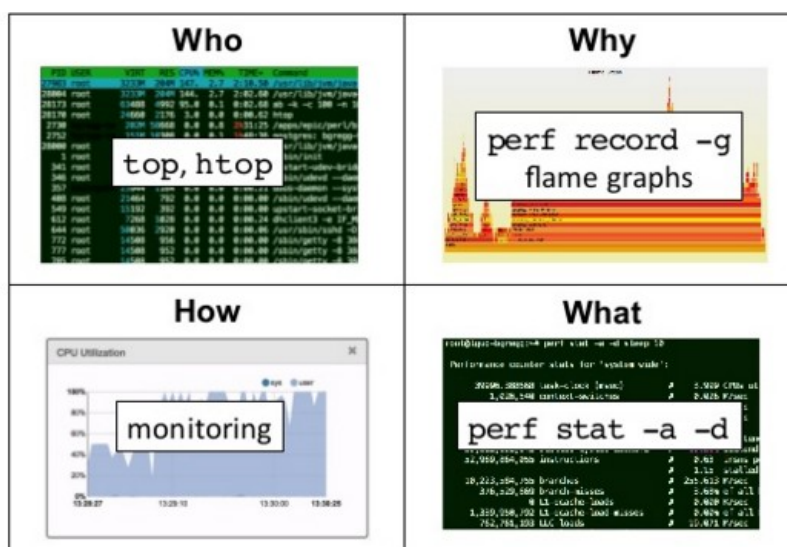
The “Who” is usually easy: tools like top, htop, mpstat, perf top, etc will show you.

The “How” is also relatively easy- lots of system monitoring tools available (sar, nagios, cacti, nmon, sysmon, etc).

The problem areas are usually the “Why” and “What”!

- Who
- Why
- What
- How

CPU Tools



CPU Profiling with Flame Graphs

<http://www.brendangregg.com/flamegraphs.html>

- In a “regular” analysis with perf, you'd sample cpu (perf record ...) and then report your findings (perf report ...)
- Here's a quick sample:

- Sampling full stack traces at 99 Hertz:

```
# perf record -F 99 -ag -- sleep 30
[ perf record: Woken up 9 times to write data ]
[ perf record: Captured and wrote 2.745 MB perf.data (~119930 samples) ]
# perf report -n --stdio
1.40% 162      java [kernel.kallsyms]          [k] _raw_spin_lock
      |
      |-- _raw_spin_lock
      |   |--63.21%-- try_to_wake_up
      |       |--63.91%-- default_wake_function
      |           |--56.11%-- _wake_up_common
      |               _wake_up_locked
      |               ep_poll_callback
      |               _wake_up_common
      |               wake_up_sync_key
      |           |--59.19%-- sock_def_readable
      |
      [...78,000 lines truncated...]
```

- Trouble is, on a production server, the output of perf report --stdio could go into tens of thousands of lines! (78k above).
- Around 8000 lines would look like this:

perf report

The image shows a dense, repetitive output of the 'perf report' command. The text is highly compressed and repetitive, showing many lines of stack trace information that are truncated or repeated, illustrating the volume of data generated in a production environment.


```
git clone --depth 1 https://github.com/brendangregg/FlameGraph
cd FlameGraph
perf record -F 99 -a -g -- sleep 30
perf script | ./stackcollapse-perf.pl | ./flamegraph.pl > perf.svg
```

- Flame Graphs:
 - **x-axis**: alphabetical stack sort, to maximize merging
 - **y-axis**: stack depth
 - **color**: random, or hue can be a dimension
 - e.g., software type, or difference between two profiles for non-regression testing ("differential flame graphs")
 - interpretation: top edge is on-CPU, beneath it is ancestry
- Just a Perl program to convert perf stacks into SVG
 - Includes JavaScript: open in a browser for interactivity
- Easy to get working

The following pages (or posts) introduce different types of flame graphs:

- [CPU](#)
- [Memory](#)
- [Off-CPU](#)
- [Hot/Cold](#)
- [Differential](#)

The example on the right is a portion of a CPU flame graph, showing MySQL codepaths that are consuming CPU cycles, and by how much.

Summary

The **x-axis shows the stack profile population, sorted alphabetically** (it is **not** the passage of time), and the **y-axis shows stack depth**. Each **rectangle** represents a **stack frame**. The wider a frame is, the more often it was present in the stacks. The **top edge shows what is on-CPU**, and beneath it is its ancestry. The colors are usually not significant, picked randomly to differentiate frames.

This visualization is fully introduced and explained in the [CPU Flame Graphs](#) page, and in the [presentation](#) below.

...

Generating a FlameGraph

1. Install the FlameGraph scripts [from here](#) or do (in an empty folder):

```
git clone --depth 1 https://github.com/brendangregg/FlameGraph.git
```

2. Generate a perf.data sample set

```
perf record -F 99 --call-graph dwarf [-a]|[-p pid]
```

Generates perf.data

[-a => all cpu's; in effect, if specified, the sample is system-wide

-p => sample a particular process.]

3. Convert using perf script:

```
perf script > perfscripout.dat
```

Generate the “flamegraph” SVG

```
cat perfscripout.dat | FlameGraph/stackcollapse-perf.pl |  
    FlameGraph/flamegraph.pl > out.svg
```

4. Open the SVG in a web browser, move the mouse over stack frames (**each rectangle** represents a single stack frame: the **width** is representative of the **frequency** of the function call, the **height** is representative of the **depth** of the stack; the order from left-to-right is just **alphabetical**, it's *not* a timeline), click on stack frames to zoom into it.

In effect: **the hottest code-paths are the widest sections!**

The top-edge is the stuff on-CPU; beneath is ancestry.

To zoom out to the whole picture again, click on the “Reset Zoom” hyperlink on the upper-left corner.

Enjoy!

Participants with the [Seawolf pendrive](#) :

Prebuilt scripts will let you generate FlameGraph's very easily.

Just navigate to the ~/kaiwanTECH/FlameGraph folder, read the Readme_FlameGraph.txt there and proceed.

Source: CPU Flame Graphs

...

Description

I'll explain this carefully: it may look similar to other visualizations from profilers, but it is different.

- Each box represents a function in the stack (a "stack frame").
- The **y-axis** shows stack depth (number of frames on the stack). The top box shows the function that was on-CPU. Everything beneath that is ancestry. The function beneath a function is its parent, just like the stack traces shown earlier.
- The **x-axis** spans the sample population. It does *not* show the passing of time from left to right, as most graphs do. The left to right ordering has no meaning (it's sorted alphabetically).
- The width of the box shows the *total* time it was on-CPU or part of an ancestry that was on-CPU (based on sample count). Wider box functions may be slower than narrow box functions, or, they may simply be called more often. The call count is not shown (or known via sampling).
- The sample count can exceed elapsed time if multiple threads were running and sampled concurrently.

The colors aren't significant, and are usually picked at random to be warm colors (other meaningful palettes are supported). It's called "flame graph" as it's showing what is hot on-CPU. And, it's interactive: mouse over the SVGs to reveal details, and click to zoom.

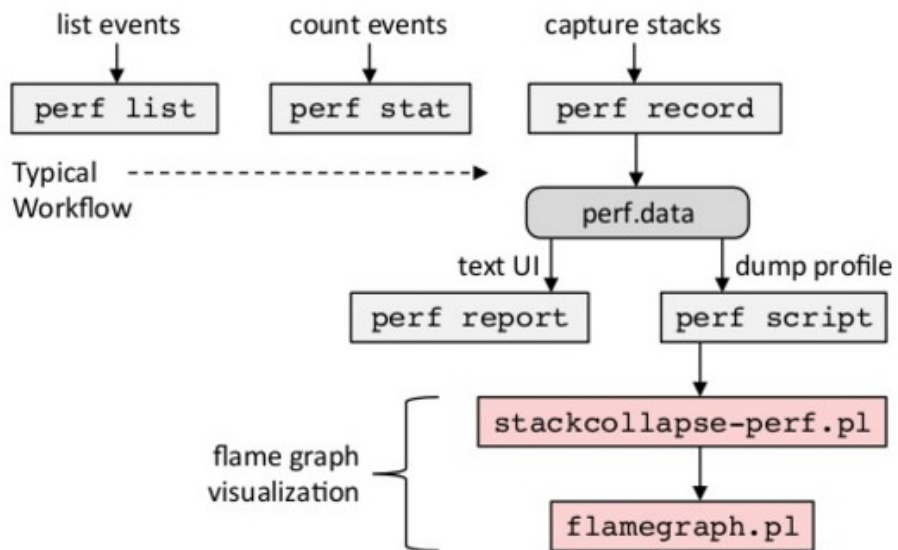
...

```
# cd FlameGraph
# perf record -F 99 -a -g -- sleep 60
# perf script | ./stackcollapse-perf.pl > out.perf-folded
# ./flamegraph.pl out.perf-folded > perf-kernel.svg
```

The `perf record` command samples at 99 Hertz (-F 99) across all CPUs (-a), capturing stack traces so that a call graph (-g) of function ancestry can be generated later. The samples are saved in a `perf.data` file, which are read by `perf script`.

...

perf Actions: Workflow



Gotchas!
(see below)


When you try to use perf

- Stacks don't work (missing)
- Symbols don't work (hex numbers)
- Can't profile Java
- Can't profile Node.js/io.js
- PMCs aren't available
- Dynamic tracing function arguments don't work
- perf locks up

PMC = (cpu) Performance Measurement Counters

The GOTCHA's!

How to *really* get started



This is my actual checklist.

IPC: (cpu) Instructions Per Cycle

Gotcha #1 Broken Stacks

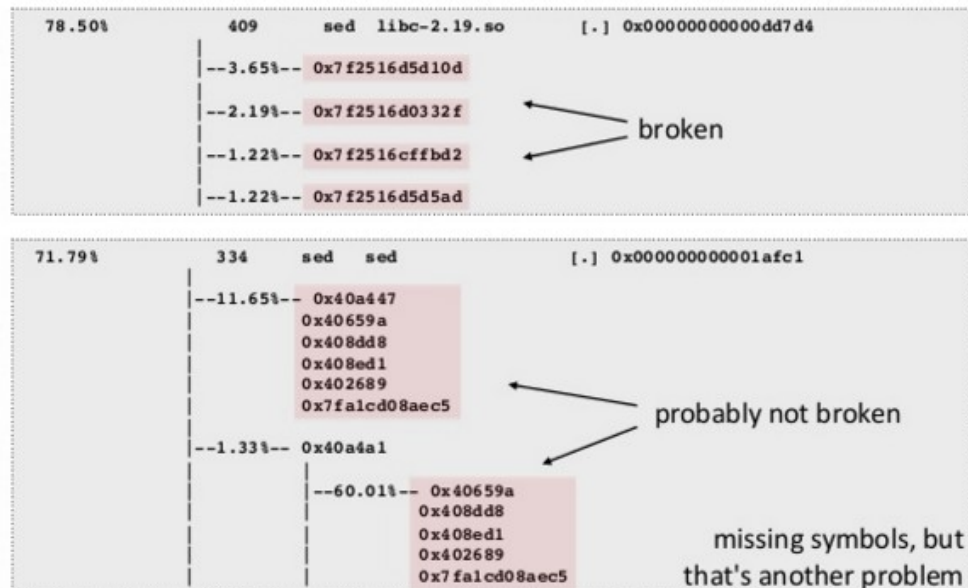
```
perf record -F 99 -a -g -- sleep 30
perf report -n --stdio
```

Start by testing stacks:

1. Take a CPU profile
2. Run perf report
3. If stacks are often < 3 frames, or don't reach "thread start" or "main", they are probably broken. Fix them.

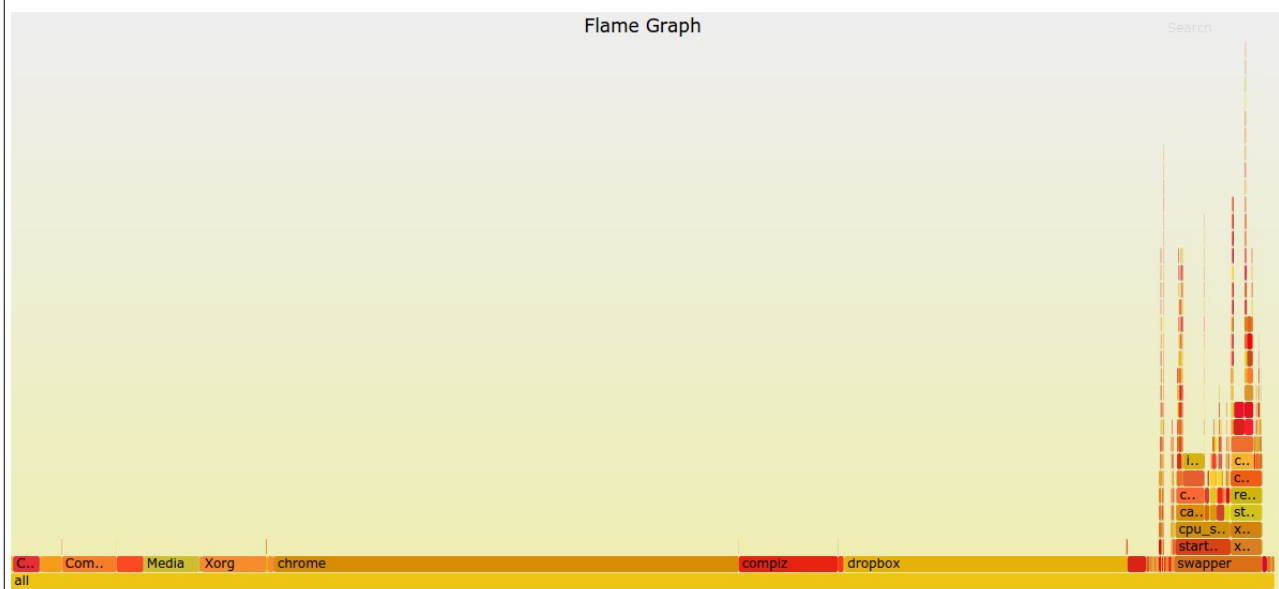
Identifying Broken Stacks (1)

Identifying Broken Stacks



Identifying Broken Stacks (2)

Symbols missing here..



Identifying Broken Stacks (3)

Notice how “low & flat” the Flame Graph shown above is: implies no stack info – broken stacks. Only on the extreme right is there a “tower” implying availability of stack info.

Turns out this was kernel-mode code.

Solution to Gotcha #1 : Broken Stacks

- Compilable code: use the '-fno-omit-frame-pointer' gcc switch *[preferred]*
- perf ... --call-graph dwarf *[use DWARF & libunwind]*
- custom stack walker (probably needs kernel support, slow)

Gotcha #2 Missing Symbols

- Missing symbols should be obvious in perf report/script:

```

71.79%          334      sed      sed          [.] 0x0000000000001afc1
|
|--11.65%-- 0x40a447
              0x40659a
              0x408dd8
              0x408ed1 ← broken
              0x402689
              0x7fa1cd08aec5

```

```

12.06%      62      sed      sed      [...] re_search_internal
|
--- re_search_internal
|
|--96.78%-- re_search_stub
|           rpl_re_search
|           match_regex
|           do_subst
|           execute_program
|           process_files
|           main
|           libc start main

```

← not broken

Fixing Symbols

- For installed packages:
 - A. Add a `-dbgsym` package, if available
 - B. Recompile from source
- For JIT (Java, Node.js, ...):
 - A. Create a `/tmp/perf-PID.map` file. `perf` already looks for this
 - B. Or use one of the new symbol loggers (see `lkml`)

```
# perf script
Failed to open /tmp/perf-8131.map, continuing without symbols
[...]
```

```
java 8131 cpu-clock:
    7fff76f2dce1 [unknown] ([vdso])
    7fd3173f7a93 os::javaTimeMillis() (/usr/lib/jvm...
    7fd301861e46 [unknown] (/tmp/perf-8131.map)
[...]
```

Compiled code: Use `gcc -g -ggdb` switches to include symbolic info
 -dbgsym package for the app / lib (even a debug *vmlinux* kernel image should be available for most platforms).

FYI: Java

perf JIT symbols: Java, Node.js

- Using the /tmp map file for symbol translation:
 - Pros: simple, can be low overhead (snapshot on demand)
 - Cons: stale symbols
 - Map format is "START SIZE symbolname"
 - Another gotcha: If perf is run as root, chown root <mapfile>
- Java
 - <https://github.com/jrudolph/perf-map-agent>
 - Agent attaches and writes the map file on demand (previous versions attached on Java start, and wrote continually)
- Node.js
 - `node --perf_basic_prof` writes the map file continually
 - Available from 0.11.13+
 - Currently has a file growth issue; see my patch in <https://code.google.com/p/v8/issues/detail?id=3453>

Gotcha #3 Guest PMCs

- Using PMCs from a Xen guest (currently):

```
# perf stat -a -d sleep 5

Performance counter stats for 'system wide':

   10003.718595 task-clock (msec)    #    2.000 CPUs utilized    [100.00%]
        323 context-switches       #   0.032 K/sec             [100.00%]
         17 cpu-migrations          #   0.002 K/sec             [100.00%]
        233 page-faults             #   0.023 K/sec
<not supported> cycles
<not supported> stalled-cycles-frontend
<not supported> stalled-cycles-backend
<not supported> instructions
<not supported> branches
<not supported> branch-misses
<not supported> Ll-dcache-loads
<not supported> Ll-dcache-load-misses
<not supported> LLC-loads
<not supported> LLC-load-misses

   5.001607197 seconds time elapsed
```


<p><i>Virtual Machines</i>: they could, but most do not as of today, implement CPU PMCs (Performance Measurement Counters). Minimally require: IPC (Instructions Per Cycle) Workaround: making use of MSRs (Model Specific Registers): usually are implemented in VMs, and can help.</p>

More examples explained

...

(More) Usage Examples

These example sequences have been chosen to illustrate some different ways that **perf** is used, from gathering to reporting.

Performance counter summaries, including IPC, for the **gzip** command:

```
# perf stat gzip largefile
```

Count all scheduler process events for 5 seconds, and count by tracepoint:

```
# perf stat -e 'sched:sched_process_*' -a sleep 5
```

Trace all scheduler process events for 5 seconds, and count by both tracepoint and process name:

```
# perf record -e 'sched:sched_process_*' -a sleep 5
# perf report
```

Trace all scheduler process events for 5 seconds, and dump per-event details:

```
# perf record -e 'sched:sched_process_*' -a sleep 5
# perf script
```

Trace **read()** syscalls, when requested bytes is less than 10:

```
# perf record -e 'syscalls:sys_enter_read' --filter 'count < 10' -a
```

Sample CPU stacks at 99 Hertz, for 5 seconds:

```
# perf record -F 99 -ag -- sleep 5
# perf report
```

Dynamically instrument the kernel `tcp_sendmsg()` function, and trace it for 5 seconds, with stack traces:

```
# perf probe --add tcp_sendmsg
# perf record -e probe:tcp_sendmsg -ag -- sleep 5
# perf probe --del tcp_sendmsg
# perf report
```

Deleting the tracepoint (`--del`) wasn't necessary; I included it to show how to return the system to its original state.

...

Profiling vs Tracing

- Profiling takes samples. Tracing records every event.
- There are many tracers for Linux (SystemTap, ktap, etc), but only two in mainline: `perf_events` and `ftrace`

Linux Tracing Stack

one-liners:	<i>many</i>
front-end tools:	<code>perf</code> , <code>perf-tools</code>
tracing frameworks:	<code>perf_events</code> , <code>ftrace</code> , <code>eBPF</code> , ...
tracing instrumentation:	tracepoints, kprobes, uprobes

Tracing Example

```
# perf record -e block:block_rq_insert -a
^C[ perf record: Woken up 1 times to write data ]
[ perf record: Captured and wrote 0.172 MB perf.data (~7527 samples) ]

# perf script
[...]
```

java	9940	[015]	1199510.044783:	block_rq_insert:	202,1 R 0 ()	4783360 + 88	[java]
java	9940	[015]	1199510.044786:	block_rq_insert:	202,1 R 0 ()	4783448 + 88	[java]
java	9940	[015]	1199510.044786:	block_rq_insert:	202,1 R 0 ()	4783536 + 24	[java]
java	9940	[000]	1199510.065195:	block_rq_insert:	202,1 R 0 ()	4864088 + 88	[java]

[...] ↑

process PID [CPU] timestamp: eventname: format string

```
include/trace/events/block.h: java 9940 [015] 1199510.044783: block_rq_insert: 202,1 R 0 ()
4783360 + 88 [java]
DECLARE_EVENT_CLASS(block_rq,
[...]
```

```
TP_printk("%d,%d %s %u (%s) %llu + %u [%s]",
          MAJOR(__entry->dev), MINOR(__entry->dev),
          __entry->rwbs, __entry->bytes, __get_str(cmd),
          (unsigned long long)__entry->sector,
          __entry->nr_sector, __entry->comm)
```

kernel source
may be the
only docs

More Examples

One-Liners: Static Tracing

```
# Trace new processes, until Ctrl-C:
perf record -e sched:sched_process_exec -a

# Trace all context-switches with stack traces, for 1 second:
perf record -e context-switches -ag -- sleep 1

# Trace CPU migrations, for 10 seconds:
perf record -e migrations -a -- sleep 10

# Trace all connect()s with stack traces (outbound connections), until Ctrl-C:
perf record -e syscalls:sys_enter_connect -ag

# Trace all block device (disk I/O) requests with stack traces, until Ctrl-C:
perf record -e block:block_rq_insert -ag

# Trace all block device issues and completions (has timestamps), until Ctrl-C:
perf record -e block:block_rq_issue -e block:block_rq_complete -a

# Trace all block completions, of size at least 100 Kbytes, until Ctrl-C:
perf record -e block:block_rq_complete --filter 'nr_sector > 200'

# Trace all block completions, synchronous writes only, until Ctrl-C:
perf record -e block:block_rq_complete --filter 'rwbs == "WS"'

# Trace all block completions, all types of writes, until Ctrl-C:
perf record -e block:block_rq_complete --filter 'rwbs ~ "*W*"'

# Trace all ext4 calls, and write to a non-ext4 location, until Ctrl-C:
perf record -e 'ext4:*' -o /tmp/perf.data -a
```

Perf – Dynamic Tracing

One can add tracepoints dynamically via perf probe!

This works on any functions within the kernel, except for those marked as 'static'.

Even works for functions that are private (that are not exported via EXPORT_SYMBOL macro).

One-Liners: Dynamic Tracing

```
# Add a tracepoint for the kernel tcp_sendmsg() function entry (--add optional):
perf probe --add tcp_sendmsg

# Remove the tcp_sendmsg() tracepoint (or use --del):
perf probe -d tcp_sendmsg

# Add a tracepoint for the kernel tcp_sendmsg() function return:
perf probe 'tcp_sendmsg%return'

# Show avail vars for the tcp_sendmsg(), plus external vars (needs debuginfo):
perf probe -V tcp_sendmsg --externs

# Show available line probes for tcp_sendmsg() (needs debuginfo):
perf probe -L tcp_sendmsg

# Add a tracepoint for tcp_sendmsg() line 81 with local var seglen (debuginfo):
perf probe 'tcp_sendmsg:81 seglen'

# Add a tracepoint for do_sys_open() with the filename as a string (debuginfo):
perf probe 'do_sys_open filename:string'

# Add a tracepoint for myfunc() return, and include the retval as a string:
perf probe 'myfunc%return +0($retval):string'

# Add a tracepoint for the user-level malloc() function from libc:
perf probe -x /lib64/libc.so.6 malloc

# List currently available dynamic probes:
perf probe -l
```

Eg.1 : Kernel function: ip_rcv

```
# perf probe --add ip_rcv
```

Added new event:

```
probe:ip_rcv          (on ip_rcv)
```

You can now use it in all perf tools, such as:

```
perf record -e probe:ip_rcv -aR sleep 1
```

```
# perf probe -l
```

```
probe:ip_rcv          (on ip_rcv)
```

```
#
```

Eg.2 : Libc function: malloc

```
# perf probe -x /lib/x86_64-linux-gnu/libc-2.21.so malloc
<< -x : -x, --exec=PATH
    Specify path to the executable or shared library file for user
space tracing.
    Can also be used with --funcs option.
```

```
>>
```

```
Added new events:
```

```
probe_libc:malloc      (on malloc in /lib/x86_64-linux-gnu/libc-2.21.so)
probe_libc:malloc_1    (on malloc in /lib/x86_64-linux-gnu/libc-2.21.so)
probe_libc:malloc_2    (on malloc in /lib/x86_64-linux-gnu/libc-2.21.so)
probe_libc:malloc_3    (on malloc in /lib/x86_64-linux-gnu/libc-2.21.so)
probe_libc:malloc_4    (on malloc in /lib/x86_64-linux-gnu/libc-2.21.so)
probe_libc:malloc_5    (on malloc in /lib/x86_64-linux-gnu/libc-2.21.so)
probe_libc:malloc_6    (on malloc in /lib/x86_64-linux-gnu/libc-2.21.so)
probe_libc:malloc_7    (on malloc in /lib/x86_64-linux-gnu/libc-2.21.so)
probe_libc:malloc_8    (on malloc in /lib/x86_64-linux-gnu/libc-2.21.so)
probe_libc:malloc_9    (on malloc in /lib/x86_64-linux-gnu/libc-2.21.so)
probe_libc:malloc_10   (on malloc in /lib/x86_64-linux-gnu/libc-2.21.so)
probe_libc:malloc_11   (on malloc in /lib/x86_64-linux-gnu/libc-2.21.so)
probe_libc:malloc_12   (on malloc in /lib/x86_64-linux-gnu/libc-2.21.so)
probe_libc:malloc_13   (on malloc in /lib/x86_64-linux-gnu/libc-2.21.so)
probe_libc:malloc_14   (on malloc in /lib/x86_64-linux-gnu/libc-2.21.so)
```

You can now use it in all perf tools, such as:

```
perf record -e probe_libc:malloc_14 -aR sleep 1
```

```
#
```

```
# perf record -e probe_libc:malloc_1 -aR --call-graph dwarf sleep 3
```

```
[ perf record: Woken up 2443 times to write data ]
```

```
Warning:
```

```
Processed 188146 events and lost 945 chunks!
```

Check IO/CPU overload!

```
[ perf record: Captured and wrote 1332.009 MB perf.data (164998 samples) ]
```

```
# 1 perf.data
```

```
-rw----- 1 root root 1.4G Dec 23 11:55 perf.data
```

```
#
```

perf report

Samples: 164K of event 'probe_libc:malloc_1', Event count (approx.): 164998

Children	Self	Command	Shared Object	Symbol
-	83.37%	83.37%	dropbox	libc-2.21.so
-				[.] malloc
-				malloc
+	42.42%		0x53a707	
+	33.49%		realloc	
+	8.44%		PyList_New	
+	6.46%		0x4c8f99	
+	3.20%		_nl_make_l10nflist	
+	1.96%		PyString_FromStringAndSize	
+	1.43%		0x45e921	
-	0.79%		CRYPTO_malloc	
-			EVP_MD_CTX_copy_ex	
+	50.05%		0x4ff3f5	
+	49.95%		0x4ff82c	
+	82.54%	0.00%	dropbox	[.] PyEval_EvalCodeEx
+	82.54%	0.00%	dropbox	[.] PyObject_Call
+	82.54%	0.00%	dropbox	[.] PyEval_EvalFrameEx
+	76.73%	0.00%	dropbox	[.] 0x000000000001e15c5
+	58.16%	0.00%	dropbox	[.] 0x00000000000206a5f
+	41.75%	0.00%	dropbox	[.] 0x00000000000206990
+	40.70%	0.00%	dropbox	[.] 0x000000000001f985a
+	35.37%	0.00%	dropbox	[.] 0x0000000000010f1b9
+	35.37%	0.00%	dropbox	[.] 0x0000000000013a707
+	34.74%	0.00%	dropbox	[.] 0x00000000000001971
+	27.92%	0.00%	dropbox	[.] realloc
+	26.55%	0.00%	dropbox	[.] 0x00000000000080a2a
+	26.17%	0.00%	dropbox	[.] 0x00000000000187c50
+	25.83%	0.00%	dropbox	[.] PyObject_CallFunctionObjArgs

Tip:

perf report -stdio

to see everything at once with expanded stack traces.

Brendan Gregg's perf-tools-unstable

Interesting:

Brendan Gregg has developed **useful and powerful bash script wrappers over ftrace, kprobes**, etc. Install the '**perf-tools-unstable**' package to try them out (*pre-installed on the [Seawolf](#) virtual appliance*).

They come with man pages too! :-)

[Git repo](#).

```
$ dpkg -L perf-tools-unstable
```

```
/.
/usr
/usr/bin
/usr/bin/bitesize
/usr/bin/cachestat
/usr/bin/execsnoop
/usr/bin/funccount
/usr/bin/funcgraph
```

```

/usr/bin/funclower
/usr/bin/functrace
/usr/bin/iolatency
/usr/bin/iosnoop
/usr/bin/killsnoop
/usr/bin/kprobe
/usr/bin/opensnoop
/usr/bin/perf-stat-hist
/usr/bin/reset-ftrace
/usr/bin/syscount
/usr/bin/tcpretrans
/usr/bin/tpoint
/usr/bin/uprobe
...
$

```

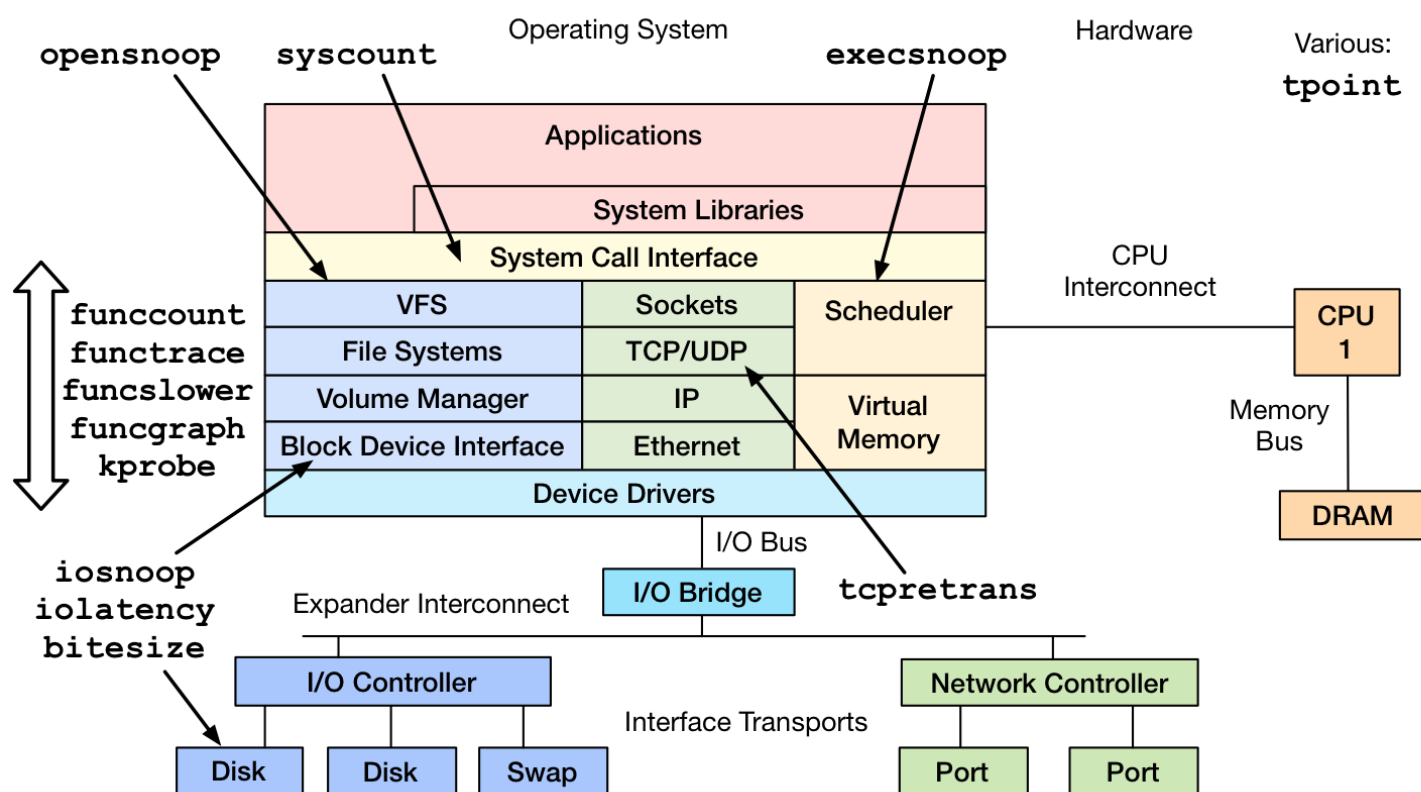
(On some installations, the script name is suffixed with -perf).

***!WARNING! Many of these tools are considered to be Experimental and Unstable!
Should only be used on kernel ver >= 4.0.***

[Source](#)

Linux Performance Observability Tools: Perf Tools, Brendan Gregg

Linux Performance Observability Tools: perf-tools



!WARNING! Many of these tools are considered to be Experimental and Unstable!

DTrace for Linux 2016, Oct 2016

or

BPF (Berkeley Packet Filter) Tracing

Linux 4.9 on now has advanced tracing capabilities, similar to Solaris's dtrace.

“... On Linux, you can now analyze the performance of applications and the kernel using production-safe low-overhead custom tracing, with latency histograms, frequency counts, and more.”

- eBPF (enhanced Berkeley Packet Filter) timed sampling now merged into 4.9-rc1.

Brendan Gregg

Talks

Velocity 2015

At Velocity 2015, I gave a 90 minute tutorial on Linux performance tools, summarizing performance observability, benchmarking, tuning, static performance tuning, and tracing tools. I also covered performance methodology, and included some live demos. This should be useful for everyone working on Linux systems.

A video of the talk is on youtube ([playlist](#); [part 1](#), [part 2](#)) and the slides are on [slideshare](#) or as a [PDF](#).

This was similar to my [SCaLE11x](#) and [LinuxCon](#) talks, however, with 90 minutes I was able to cover more tools and methodologies, making it the most complete tour of the topic I've done. I also posted about it on the [Netflix Tech Blog](#).

SCALE13x (2015)

At the Southern California Linux Expo ([SCALE 13x](#)), I << **Brendan Gregg** >> gave a talk on **Linux Profiling at Netflix using perf_events (aka "perf")**, covering CPU profiling and a tour of other features. This talk includes a crash course on perf_events that you may find useful, plus it covers gotchas such as fixing stack traces and symbols when profiling Java and Node.js.

A video of the talk is on [youtube](#), and the slides are on [slideshare](#).

In a [post](#) about this talk, I included the interactive **CPU flame graph SVG** I was demonstrating.


...

<<

Participants are *strongly encouraged* to see Brendan Gregg's video's!
The SCALE13x video focusses on perf.

>>

Linux Operating System Specialized


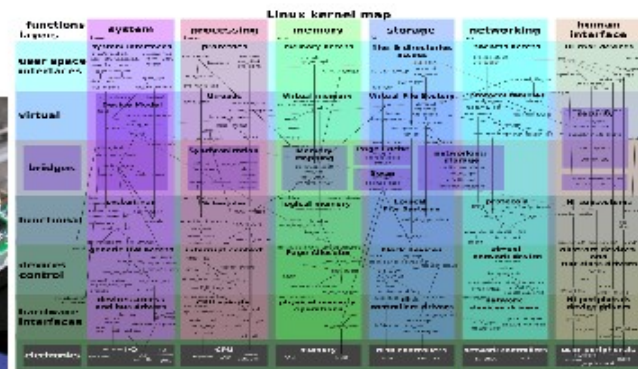


The highest quality Training on:


Linux Fundamentals, CLI and Scripting
Linux Systems Programming
Linux Kernel Internals
Linux Device Drivers
Embedded Linux
Linux Debugging Techniques
New! *Linux OS for Technical Managers*

Please do visit our website for details:

<http://kaiwantech.in>



Linux Operating System Specialized






The highest quality Training on:

Linux Fundamentals, CLI and Scripting
Linux Systems Programming
Linux Kernel Internals
Linux Device Drivers
Embedded Linux
Linux Debugging Techniques
New! *Linux OS for Technical Managers*

Please do visit our website for details:

<http://kaiwantech.in>



<http://kaiwantech.in>

kaiwanTECH Linux OS Corporate Training Programs

Please do check out our current offering of world-class, seriously-valuable, high on returns, technical Linux OS corporate training programs here: <http://bit.ly/ktcorp>