



Linux System Monitoring

Important Notice : Courseware - Legal

This courseware is both the product of the author and of freely available opensource and/or public domain materials. Wherever external material has been shown, it's source and ownership have been clearly attributed. We acknowledge all copyrights and trademarks of the respective owners.

The contents of the **courseware PDFs are considered proprietary** and thus cannot be copied or reproduced in any form whatsoever without the explicit written consent of the author.

Only the programs - **source code** and binaries (where applicable) - that form part of this courseware, and that are made available to the participant, are released under the terms of the [permissive MIT license](#).

Under the terms of the MIT License, you can certainly use the source code provided here; you must just attribute the original source (author of this courseware and/or other copyright/trademark holders).

VERY IMPORTANT :: Before using this source(s) in your project(s), you ***MUST*** check with your organization's legal staff that it is appropriate to do so.

The courseware PDFs are ***not*** under the MIT License, they are to be kept confidential, non-distributable without consent, for your private internal use only.

The duration, contents, content matter, programs, etc. contained in this courseware and companion participant VM are subject to change at any point in time without prior notice to individual participants.

Care has been taken in the preparation of this material, but there is no warranty, expressed or implied of any kind, and we can assume no responsibility for any errors or omissions. No liability is assumed for incidental or consequential damages in connection with or arising out of the use of the information or programs contained herein.

2000-2017 Kaiwan N Billimoria
kaiwanTECH, Bangalore, India.

kaiwanTECH Linux OS Corporate Training Programs

Please do check out our current offering of world-class, seriously-valuable, high on returns, technical Linux OS corporate training programs here: <http://bit.ly/ktcorp>

References:

[“Linux Performance and Tuning Guidelines”, IBM.](#)

[Linux Performance, Brendan Gregg](#)

This page links to various Linux performance material I've created, including the tools maps on the right. These show: [Linux observability tools](#), [Linux benchmarking tools](#), [Linux tuning tools](#), and [Linux sar](#). For more diagrams, see my slide decks below.

[Linux Tools for the serious Systems Programmer](#)

Saturday 28 December 2013 05:30 PM

A fairly exhaustive list, with links, of tools (and techniques) one can use for debugging on the Linux OS.

[The Different Types of Server Monitoring Software](#)

Here are the various types of server monitoring software and what they specialize in, in regards to different performance metrics.

SO: [How to determine CPU and memory consumption from inside a process?](#)

What are we really trying to achieve here?

[From Brendan Gregg's blog:](#)

“

Performance Analysis Methodology

A performance analysis methodology is a procedure that you can follow to analyze system or application performance. These **generally provide a starting point and then guidance to root cause**, or causes. Different methodologies are suited for solving different classes of issues, and you may try more than one before accomplishing your goal.

Analysis without a methodology can become a fishing expedition, where metrics are examined ad hoc, until the issue is found – if it is at all.

Methodologies documented in more detail on this site are:

- [The USE Method](#): for finding resource bottlenecks
- [The TSA Method](#): for analyzing application time
- [Off-CPU Analysis](#): for analyzing any type of thread wait latency
- [Active Benchmarking](#): for accurate and successful benchmarking.

...

USE Method

For every resource, check:

Utilization
Saturation
Errors

...

”

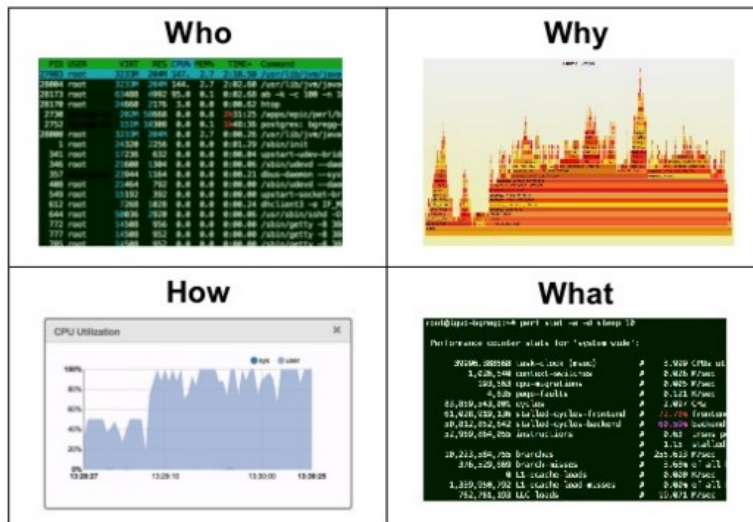
CPU Workload Characterization

[<< Source >>](#)

- Who
- Why
- What
- How

CPU Tools

Clip slide



The “Who” is usually easy: tools like top, htop, mpstat, perf top, etc will show you.
 The “How” is also relatively easy- lots of system monitoring tools available (sar, nagios, cacti, nmon, sysmon, etc).

The problem areas are usually the “Why” and “What”!

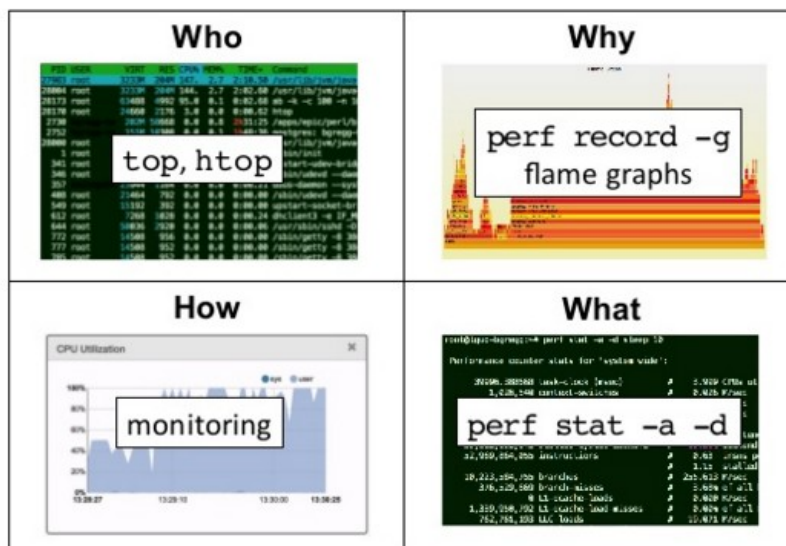
- Who
- Why
- What
- How

P.T.O. -->

- For CPUs:
 1. **Who:** which PIDs, programs, users
 2. **Why:** code paths, context
 3. **What:** CPU instructions, cycles
 4. **How:** changing over time
- Can you currently answer them? How?

Here's how:

CPU Tools



From Brendan Gregg's blog:

“Broken Linux Performance Tools (SCaLE14x, 2016)

At the Southern California Linux Expo ([SCaLE 14x](#)), I gave a talk on Broken Linux Performance Tools. This was a follow-on to my earlier Linux Performance Tools talk originally at SCaLE11x (and more recently at [Velocity](#) as a tutorial). This broken tools talk was a tour of common problems with Linux system tools, metrics, statistics, visualizations, measurement overhead, and benchmarks. It also includes advice on how to cope (the green "What You Can Do" slides). A video of the talk is on [youtube](#) and the slides are on [slideshare](#) or as a [PDF](#). ...”

Tools for Working with Processes

The participant should try out these utilities, tips, etc on their system.

ps

`man ps` [`ps(1)`] for all details of course.

Tips:

PROCESS FLAGS (usually the 1st field - 'F')

The sum of these values is displayed in the "F" column, which is provided by the flags output specifier:

- 1 forked but didn't exec
- 4 used super-user privileges

PROCESS STATE CODES (usually the 2nd field - 'STAT')

Here are the different values that the s, stat and state output specifiers (header "STAT" or "S") will display to describe the state of a process:

- D uninterruptible sleep (usually IO)
- R running or runnable (on run queue)
- S interruptible sleep (waiting for an event to complete)
- T stopped, either by a job control signal or because it is being traced
- W paging (not valid since the 2.6.xx kernel)
- X dead (should never be seen)
- Z defunct ("zombie") process, terminated but not reaped by its parent

For BSD formats and when the stat keyword is used, additional characters may be displayed:

- < high-priority (not nice to other users)
- N low-priority (nice to other users)
- L has pages locked into memory (for real-time and custom IO)
- s is a session leader
- l is multi-threaded (using CLONE_THREAD, like NPTL pthreads do)
- + is in the foreground process group

Also, regarding 'STAT' column:

```
s : is a session-leader
I : multithreaded application (uses CLONE_THREAD)
+ : is in the foreground process group
```

```
ps -l      : see details
ps -L      : see individual threads belonging to a process;
[usermode PID == kernel TGID  (the actual PID for userspace)
 usermode Thread == kernel PID ].
```

```
ps fax     : show tree
ps faxj    : show tree with more detail
```

```
ps -Aefww : show all arguments
```

ps : Customized Output Fields

To see every process with a user-defined format:

```
ps -eo pid,tid,class,rtprio,ni,pri,psr,pcpu,stat,wchan:14,comm
ps axo stat,euid,ruid,tt,tpgid,sess,pgrp,ppid,pid,pcpu,comm
ps -Ao pid,tt,user,fname,tmout,f,wchan
```

```
ps -eo pid,comm:16,stackp
```

<< *comm:16 => 16 chars width; stackp => show stack pointer (see man page for all possible fields that can be shown with ps -eo !)* >>

Print only the process IDs of syslogd:

```
ps -C syslogd -o pid=
```

Print only the name of PID 42:

```
ps -p 42 -o comm=
```

Also see and try:

pstree

pgrep, pkill

top

lsof : Lists Open Files

pidstat

See the post “[LINUX Quick Reference Cheat Sheet](#)” from my blog (covers vmstat, ps, top).

vmstat

VM mode, with repetition:

Repeat interval (sec)

```
# vmstat 5
```

procs		memory				swap		io		system			cpu			
r	b	swpd	free	buff	cache	si	so	bi	bo	in	cs	us	sy	id	wa	st
0	0	158496	325136	446320	2282396	0	0	0	7	470	1477	2	1	97	0	0
0	0	158496	323896	446332	2282788	0	0	3	78	544	1796	4	1	95	1	0


```
0 0 158496 324980 446336 2278512 0 0 0 48 523 1758 3 1 96 0 0
0 0 158492 321732 446336 2278476 6 0 10 23 513 1807 2 1 96 0 0
...
System mostly idle..
```

<<

Procs <i>r</i> : The number of runnable processes (running or waiting for run time) <i>b</i> : The number of processes in uninterruptible sleep	Swap <i>si</i> : Amount of memory swapped in from disk (/s). <i>so</i> : Amount of memory swapped to disk (/s).
Memory <i>swpd</i> : the amount of virtual memory used. <i>free</i> : the amount of idle memory. <i>buff</i> : the amount of memory used as buffers. <i>cache</i> : the amount of memory used as cache. <i>inact</i> : the amount of inactive memory. (-a option) <i>active</i> : the amount of active memory. (-a option)	IO <i>bi</i> : Blocks received from a block device (blocks/s). <i>bo</i> : Blocks sent to a block device (blocks/s).
System <i>in</i> : The number of interrupts per second, including the clock. <i>cs</i> : The number of context switches per second.	CPU <i>us sy id wa st</i> These are percentages of total CPU time. <i>us</i> : Time spent running non-kernel code. (user time, including nice time) <i>sy</i> : Time spent running kernel code. (system time) <i>id</i> : Time spent idle. Prior to Linux 2.5.41, this includes IO-wait time. <i>wa</i> : Time spent waiting for IO. Prior to Linux 2.5.41, included in idle. <i>st</i> : Time stolen from a virtual machine. Prior to Linux 2.6.11, unknown. (Grey color implies deprecated).

>>

```
...
procs -----memory----- --swap-- ----io---- -system- ----cpu-----
 r b  swpd  free  buff  cache  si  so  bi  bo  in  cs us sy id wa st
1 0 158488 226084 457832 2353500 0 0 1136 1421 1159 4010 21 3 69 6 0
1 0 158488 201120 457832 2354760 0 0 129 1808 1113 3572 26 2 71 1 0
1 0 158488 187140 457844 2356244 0 0 125 2308 958 2474 24 2 72 2 0
1 0 158488 165104 457944 2362996 0 0 375 1977 909 2476 25 1 72 2 0
1 0 158488 204496 457956 2364656 0 0 129 1864 1000 3431 23 2 66 8 0
...
System fairly busy..
```

```
# vmstat --help
```

Usage:

```
vmstat [options] [delay [count]]
```

Options:

```
-a, --active          active/inactive memory
-f, --forks           number of forks since boot
-m, --slabs           slabinfo
-n, --one-header      do not redisplay header
-s, --stats           event counter statistics
-d, --disk            disk statistics
-D, --disk-sum        summarize disk statistics
-p, --partition <dev> partition specific statistics
-S, --unit <char>    define display unit
-w, --wide            wide output

-h, --help            display this help and exit
-V, --version         output version information and exit
```

For more details see `vmstat(8)`.

#

dstat

<< *From the man page* >>

Dstat is a versatile replacement for vmstat, iostat and ifstat. Dstat overcomes some of the limitations and adds some extra features.

Dstat allows you to view all of your system resources instantly, you can eg. compare disk usage in combination with interrupts from your IDE controller, or compare the network bandwidth numbers directly with the disk throughput (in the same interval).

Dstat also cleverly gives you the most detailed information in columns and clearly indicates in what magnitude and unit the output is displayed. Less confusion, less mistakes, more efficient.

Dstat is unique in letting you aggregate block device throughput for a certain diskset or network bandwidth for a group of interfaces, ie. you can see the throughput for all the block devices that make up a single filesystem or storage system.

Dstat allows its data to be directly written to a CSV file to be imported and used by OpenOffice, Gnumeric or Excel to create graphs.

dstat

You did not select any stats, using -cdngy by default.

```

----total-cpu-usage---- -dsk/total- -net/total- ---paging-- ---system--
usr sys idl wai hiq siq| read  writ| recv  send|  in   out | int  csw
   3   1  95   1   0   0|  41k  201k|    0    0 | 119B 1163B| 508  3439
  18   3  79   0   0   0|    0    0 |1130B 210B|    0    0 | 705  1851
  24   3  73   0   0   0|    0    0 | 140B   70B|    0    0 | 755  1769
  23   3  73   1   0   0|    0  112k| 226B  140B|    0    0 | 770  1971
   1   1  87  10   0   0|    0  472k| 986B  955B|    0    0 | 654  2328
...

```

A very useful feature of dstat is the ability to quickly look up – at the granularity of a process – which process(es) is using the maximum resources.

A sample bash script function to achieve this:

```

dtop ()
{
    DLY=5;
    echo dstat --time --top-io-adv --top-cpu --top-mem ${DLY};
    dstat --time --top-io-adv --top-cpu --top-mem ${DLY}
}
...

```

\$ dtop

```

dstat --time --top-io-adv --top-cpu --top-mem 5
----system-----most-expensive-i/o-process-----most-expensive--most-expensive-
time      |process                pid read write cpu|  cpu process  | memory process

```

```

24-12 12:15:09|dropbox      5898 2697k 31k1.2%|dropbox      1.2|dropbox      706M
24-12 12:15:14|chrome      3295 8811B9546B0.2%| --enable-cra0.6|dropbox      706M
24-12 12:15:19|chrome      3295 6208B4002B0.3%| --enable-cra0.6|dropbox      706M
24-12 12:15:24|chrome      3295 5952B3379B0.2%| --enable-cra0.7|dropbox      706M
24-12 12:15:26|chrome      3295 32k4226B0.4%|cinnamon    0.5|dropbox      706M
...
$

```

htop

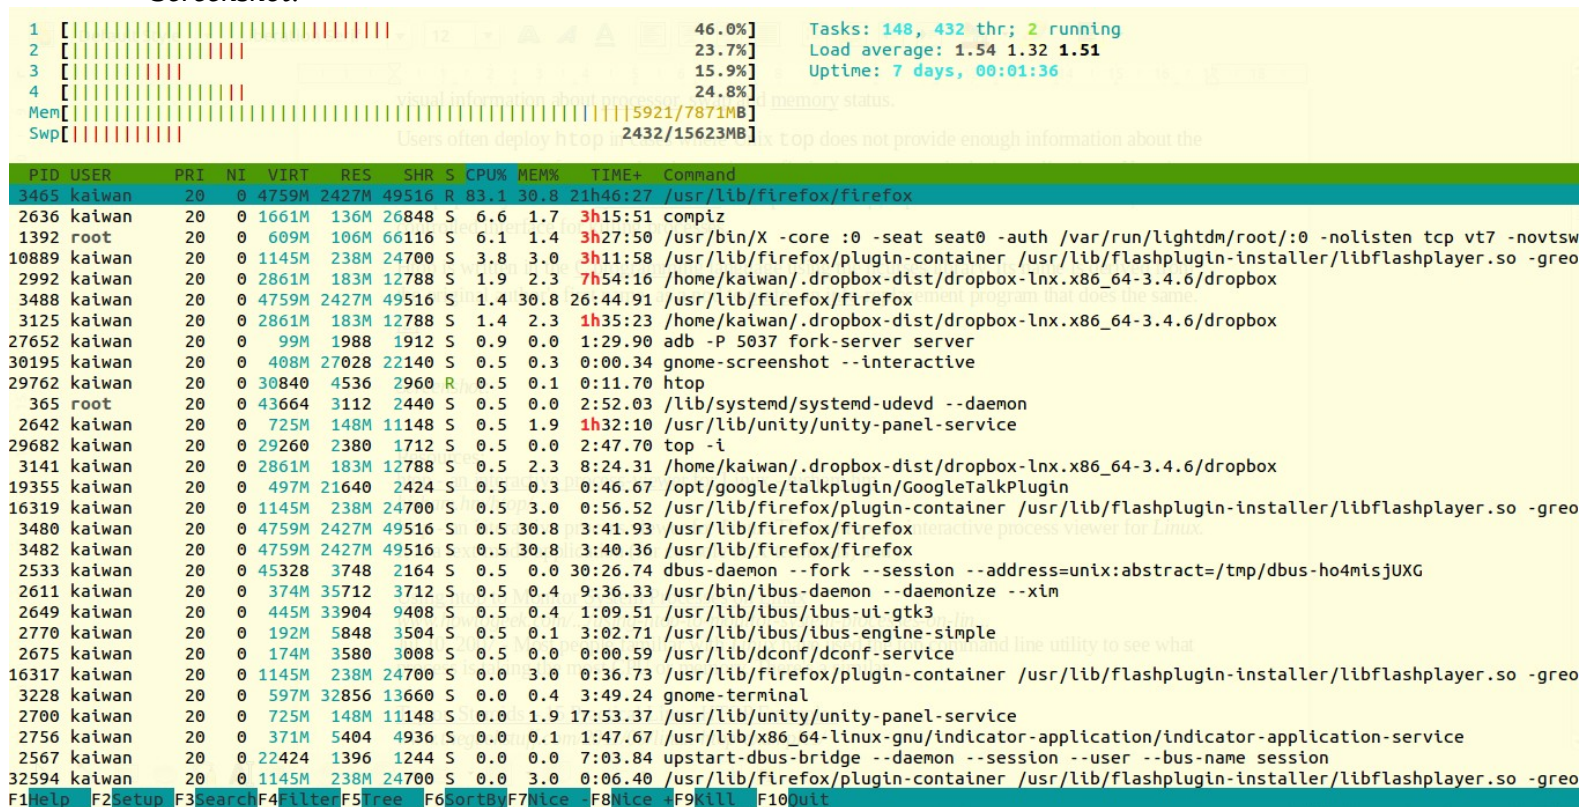
Htop is an interactive [system-monitor](#) process-viewer written for Linux. It is designed as an alternative to the Unix program [top](#). It shows a frequently updated list of the processes running on a computer, normally ordered by the amount of CPU usage. Unlike `top`, `htop` provides a full list of processes running, instead of the top resource-consuming processes. Htop uses color and gives visual information about [processor](#), [swap](#) and [memory](#) status.

Users often deploy `htop` in cases where Unix `top` does not provide enough information about the systems processes, for example when trying to find minor [memory leaks](#) in applications. Htop is also popularly used as a [system monitor](#). Compared to `top`, it provides a more convenient, cursor-controlled interface for killing processes.

Htop is written in the [C programming language](#) using the [ncurses](#) library. Its name is derived from the original author's first name, as a nod to [pinfo](#), an [info](#)-replacement program that does the same. [\[2\]](#)

P.T.O.

Screenshot:



Resources:

[htop - an interactive process-viewer for Linux - hisham.hm](http://hisham.hm/htop/)

hisham.hm/htop/

htop - an interactive process viewer for *Linux*. This is *htop*, an interactive process viewer for *Linux*. It is a text-mode application (for console or X terminals) and ...

[Using htop to Monitor System Processes on Linux](http://www.howtogeek.com/.../using-htop-to-monitor-system-processes-on-lin...)

www.howtogeek.com/.../using-htop-to-monitor-system-processes-on-lin...

Jul 10, 2007 - Most people familiar with *Linux* have used the *top* command line utility to see what process is taking the most CPU or memory. There's a similar ...

[Top on Steroids – 15 Practical Linux HTOP Examples](http://www.thegeekstuff.com/2011/09/linux-htop-examples/)

www.thegeekstuff.com/2011/09/linux-htop-examples/

Sep 14, 2011 - To use *htop*, you need to install it first. Go to *htop* download page, and download the binaries that corresponds to your *Linux* distribution and ...

[System Monitoring with htop - Rackspace](http://www.rackspace.com/.../Cloud Hosting/Cloud Servers)

[www.rackspace.com › ... › Cloud Hosting › Cloud Servers](http://www.rackspace.com/.../Cloud Hosting/Cloud Servers)

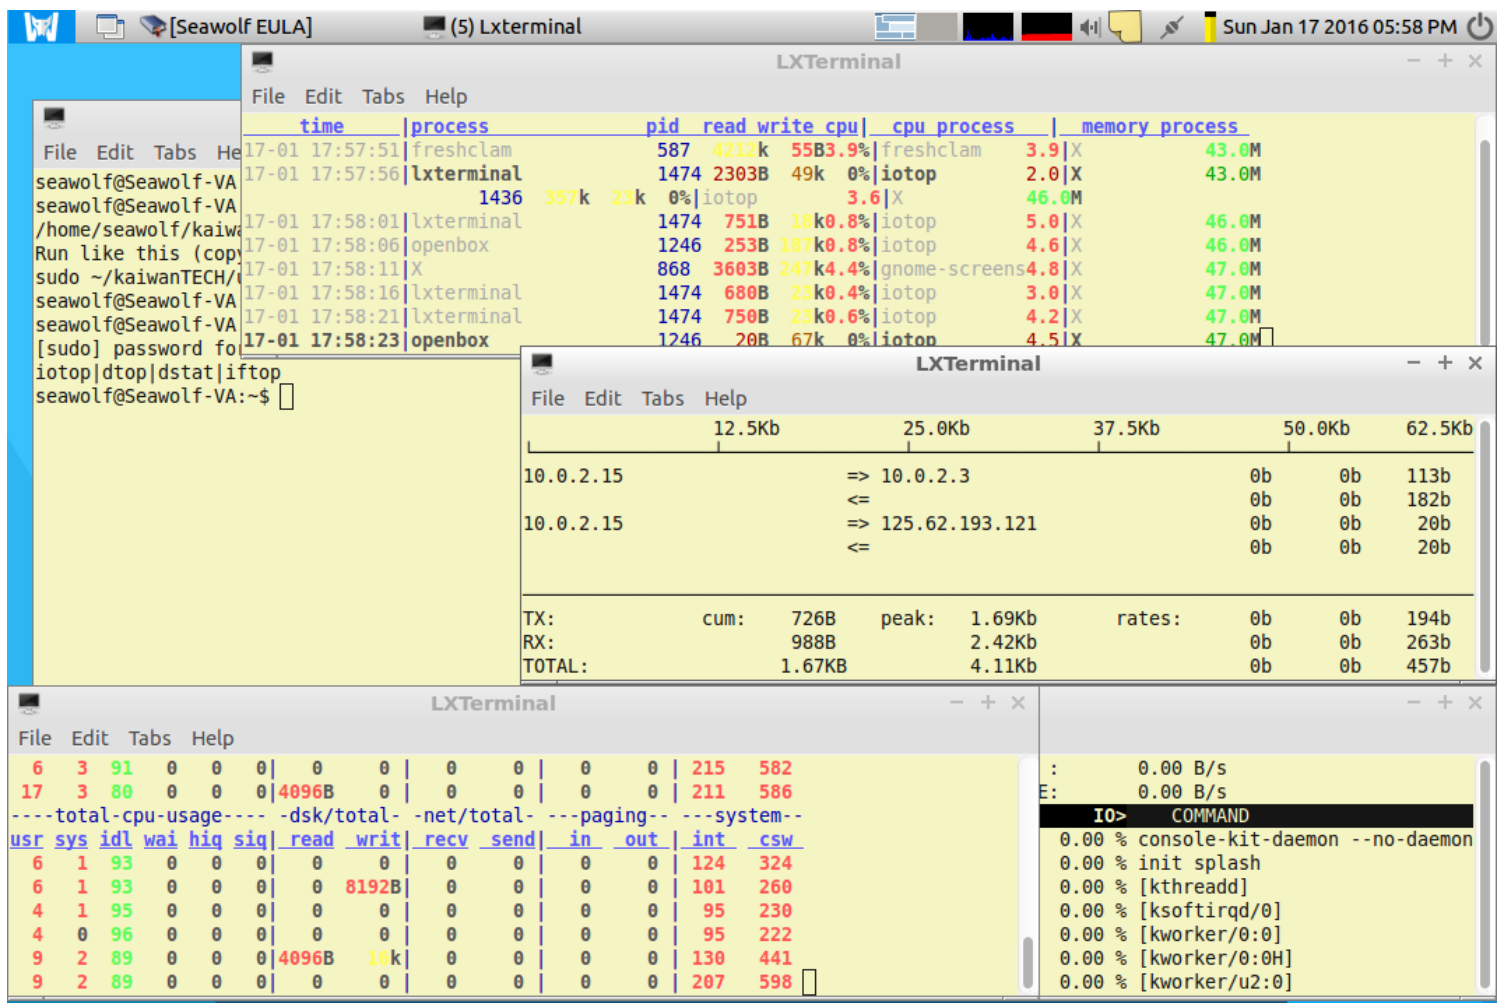
Jun 12, 2014 - *htop* is a tool in *Linux* that allows you to monitor your system's vital ... *htop* is another tool to add in your system administrator toolkit and is ...

Simple System Monitoring for a Linux Desktop

Monday 06 January 2014 12:53 PM

The Problem- What exactly is eating into my HDD / processor / network right now?? Yeah! On the (Linux) desktop, we'd like to know why things crawl along sometimes. Which process(es) is the culprit behind that disk activity, or the memory hogger, or eating up network bandwidth? Many tools exist that can help us pinpoint ... [Continue reading Simple System Monitoring for a Linux Desktop](#) →

Screenshot



Networking

Can use **iftop** to monitor networking usage.

Packet capture and analysis: [WireShark](#).

Also: [tcpdump](#), [netsniff-ng](#), [ngrep](#).


GkrellM - Graphical GUI tool for system monitoring

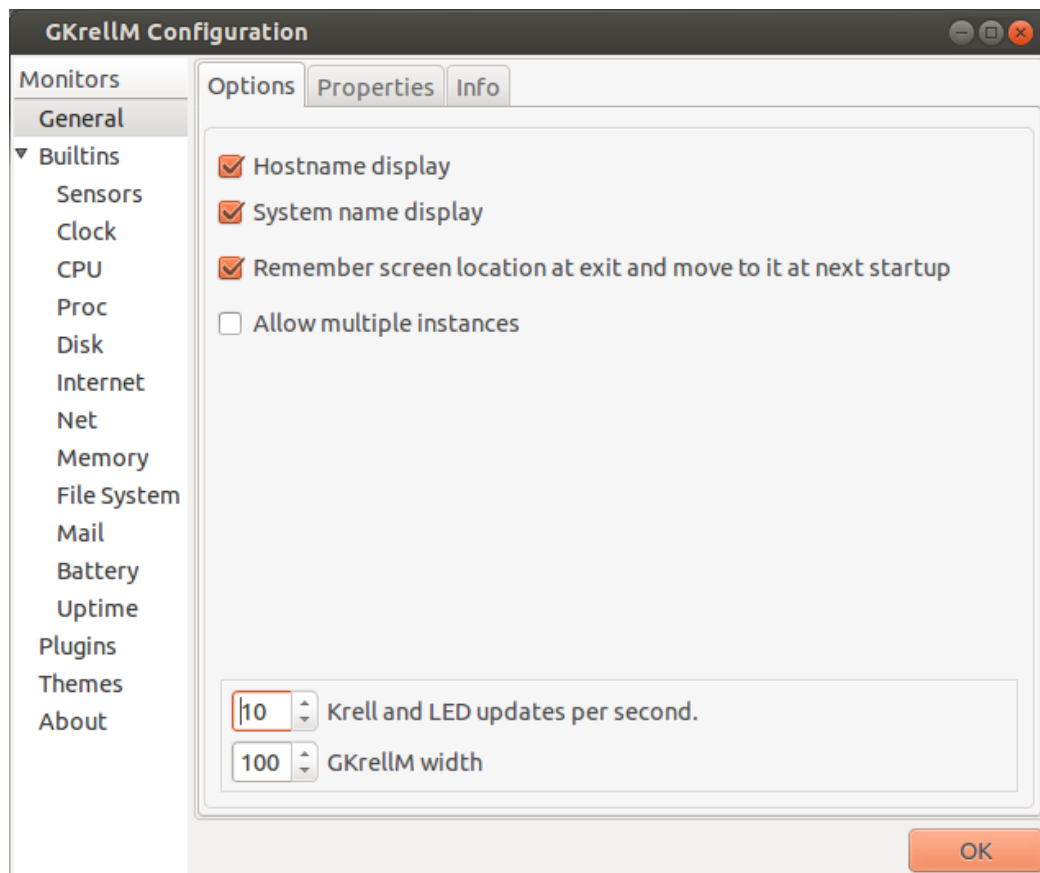
GKrellM System Monitor

Monitor for CPU, memory, disks, network, mail

With a single process, gkrellm manages multiple stacked monitors and supports applying themes to match the monitors appearance to your window manager, Gtk, or any other theme.

[Developer website](#)





nmon

nmon is a systems administrator, tuner, benchmark tool. It can display the CPU, memory, network, disks (mini graphs or numbers), file systems, NFS, top processes, resources (Linux version & processors) and on Power micro-partition information.

See the man page on nmon.

More than that:

```
$ nmon -h
```

Hint: nmon [-h] [-s <seconds>] [-c <count>] [-f -d <disks> -t -r <name>] [-x]

```
-h                FULL help information
Interactive-Mode:
read startup banner and type: "h" once it is running
For Data-Collect-Mode (-f)
-f                spreadsheet output format [note: default -s300 -c288]
optional
-s <seconds>      between refreshing the screen [default 2]
-c <number>       of refreshes [default millions]
-d <disks>        to increase the number of disks [default 256]
-t                spreadsheet includes top processes
-x                capacity planning (15 min for 1 day = -fdt -s 900 -c 96)
```

Version - nmon 14g

For Interactive-Mode

```
-s <seconds>      time between refreshing the screen [default 2]
-c <number>       of refreshes [default millions]
-g <filename>     User Defined Disk Groups [hit g to show them]
                  - file = on each line: group_name <disks list> space
```

separated

```
-                  - like: database sdb sdc sdd sde
                  - upto 64 disk groups, 512 disks per line
                  - disks can appear more than once and in many groups
-b                black and white [default is colour]
example: nmon -s 1 -c 100
```

For Data-Collect-Mode = spreadsheet format (comma separated values)

Note: use only one of f,F,z,x or X and make it the first argument

```
-f                spreadsheet output format [note: default -s300 -c288]
                  output file is <hostname>_YYYYMMDD_HHMM.nmon
-F <filename>     same as -f but user supplied filename
-r <runname>       used in the spreadsheet file [default hostname]
-t                include top processes in the output
-T                as -t plus saves command line arguments in UARG section
-s <seconds>      between snap shots
-c <number>       of snapshots before nmon stops
-d <disks>        to increase the number of disks [default 256]
-l <dpl>          disks/line default 150 to avoid spreadsheet issues. EMC=64.
-g <filename>     User Defined Disk Groups (see above) - see BBBG & DG lines
-N                include NFS Network File System
-I <percent>      Include process & disks busy threshold (default 0.1)
```

don't save or show proc/disk using less than this percent
 -m <directory> nmon changes to this directory before saving to file
 example: collect for 1 hour at 30 second intervals with top procs
 nmon -f -t -r Test1 -s30 -c120

To load into a spreadsheet:
 sort -A *nmon >stats.csv
 transfer the stats.csv file to your PC
 Start spreadsheet & then Open type=comma-separated-value ASCII file
 The nmon analyser or consolidator does not need the file sorted.

Capacity planning mode - use cron to run each day

-x sensible spreadsheet output for CP = one day
 every 15 mins for 1 day (i.e. -ft -s 900 -c 96)
 -X sensible spreadsheet output for CP = busy hour
 every 30 secs for 1 hour (i.e. -ft -s 30 -c 120)

Interactive Mode Commands

key --- Toggles to control what is displayed ---
 h = Online help information
 r = Machine type, machine name, cache details and OS version + LPAR
 c = CPU by processor stats with bar graphs
 l = long term CPU (over 75 snapshots) with bar graphs
 m = Memory stats
 L = Huge memory page stats
 V = Virtual Memory and Swap stats
 k = Kernel Internal stats
 n = Network stats and errors
 N = NFS Network File System
 d = Disk I/O Graphs
 D = Disk I/O Stats
 o = Disk I/O Map (one character per disk showing how busy it is)
 o = User Defined Disk Groups
 j = File Systems
 t = Top Process stats use 1,3,4,5 to select the data & order
 u = Top Process full command details
 v = Verbose mode - tries to make recommendations
 b = black and white mode (or use -b option)
 . = minimum mode i.e. only busy disks and processes

key --- Other Controls ---
 + = double the screen refresh time
 - = halves the screen refresh time
 q = quit (also x, e or control-C)
 0 = reset peak counts to zero (peak = ">")
 space = refresh screen now

Startup Control

If you find you always type the same toggles every time you start
 then place them in the NMON shell variable. For example:
 export NMON=cmdrvtan

Others:

- a) To you want to stop nmon - kill -USR2 <nmon-pid>
- b) Use -p and nmon outputs the background process pid
- c) To limit the processes nmon lists (online and to a file)
 Either set NMONCMD0 to NMONCMD63 to the program names
 or use -C cmd:cmd:cmd etc. example: -C ksh:vi:syncd
- d) If you want to pipe nmon output to other commands use a FIFO:

```
mkfifo /tmp/mypipe
nmon -F /tmp/mypipe &
grep /tmp/mypipe
```

e) If nmon fails please report it with:

- 1) nmon version like: 14g
- 2) the output of `cat /proc/cpuinfo`
- 3) some clue of what you were doing
- 4) I may ask you to run the debug version

Developer Nigel Griffiths

Feedback welcome - on the current release only and state exactly the problem

No warranty given or implied.

\$

```
$ export NMON=cmVn ; nmon
```

<<

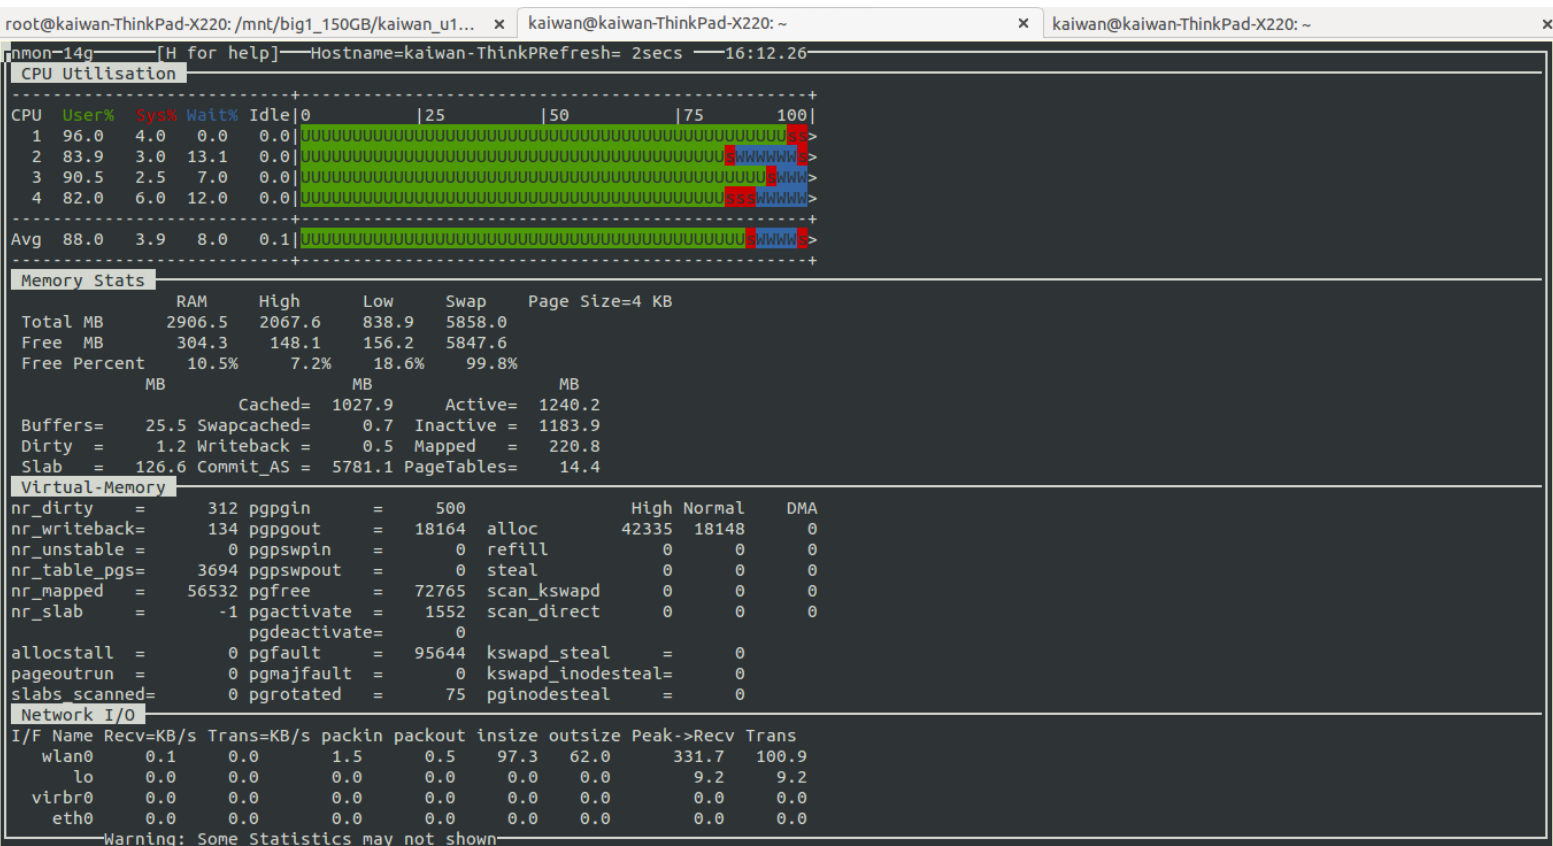
c = CPU by processor stats with bar graphs

m = Memory stats

V = Virtual Memory and Swap stats

n = Network stats and errors

>>



```
$ export NMON=duv ; nmon
```

```
<<
```

```
d   = Disk I/O Graphs
```

```
u   = Top Process full command details
```

```
v   = Verbose mode - tries to make recommendations
```

```
>>
```

```

root@kaiwan-ThinkPad-X220: /mnt/big1_150GB/kaiwan_u1... x kaiwan@kaiwan-ThinkPad-X220: ~ x kaiwan@kaiwan-ThinkPad-X220: ~
nmon-14g [H for help] Hostname=kaiwan-ThinkP Refresh= 2secs 16:15.01
Disk I/O /proc/diskstats mostly in KB/s Warning:contains duplicates
DiskName Busy Read WriteKB|0|25|50|75|100|
sda 34% 130.7 4742.1|2|
sda1 0% 0.0 0.0|>|
sda2 0% 0.0 0.0|>|
sda3 0% 0.0 0.0|>|
sda4 0% 0.0 0.0|>|
sda5 0% 0.0 0.0|>|
sda6 30% 122.8 1960.2|2|
sda7 13% 7.9 2781.9|2|
Totals Read-MB/s=0.3 Writes-MB/s=9.3 Transfers/sec=88.1
Top Processes Procs=261 mode=3 (1=Basic, 3=Perf 4=Size 5=I/O)
PID %CPU ResSize Command
Used KB
32314 51.0 42332 cc1 =506
13373 10.9 566280 /usr/lib/firefox/firefox =506
6676 8.4 88480 compiz =5for
3625 4.0 119684 /usr/bin/X -core :0 -seat seat0 -auth /var/run/lightdm/root/:0 -nolisten tcp vt7 -novtswitch =506
10161 2.0 19036 gnome-terminal =511,
6352 1.5 30736 /usr/lib/unity/unity-panel-service manag
e 17883 1.0 21620 plugin-containe g
e 30595 1.0 16584 gnome-screensho g
e 7 0.5 0 [rcu_sched] g
e 1996 0.5 8332 /usr/sbin/preload -s /var/lib/preload/preload.state g
e 5332 0.5 1900 /opt/teamviewer/tv_bin/wine/bin/wineserver g
e 6274 0.5 2352 dbus-daemon --fork --session --address=unix:abstract=/tmp/dbus-crGEVnXMyB g
e 6416 0.5 17532 /usr/lib/i386-linux-gnu/bamf/bamfdemon g
e 6697 0.5 13524 indicator-multiloat g
e 7365 0.5 132376 /home/kaiwan/.dropbox-dist/dropbox-lnx.x86-3.0.3/dropbox g
e 13483 0.5 3084 nmon r
e 18689 0.5 0 kworker/2:2 r
30948 0.5 3572 make g
e 32287 0.5 2136 make g
e 1 0.0 2504 /sbin/init ,
2 0.0 0 [kthreadd] r
3 0.0 0 [ksoftirqd/0]
5 0.0 0 [kworker/0:0H]
Warning: Some Statistics may not shown

```

Using sar

<http://perso.wanadoo.fr/sebastien.godard/>

sysstat utilities (*iostat, sar, sadf, mpstat, and sa*). The sysstat utilities are a collection of performance-monitoring tools for Linux. These tools get their system information from the *proc* file system.

sar: system activity reporter: a set of very powerful performance monitoring (perfmon) and bottleneck identification tools.

sar is like a wrapper around many existing system-gathering utilities.

“The **sysstat** package contains utilities to monitor system performance and usage activity. Sysstat contains various utilities, common to many commercial Unixes, and tools you can schedule via cron to collect and historize performance and activity data.

- [iostat\(1\)](#) reports CPU statistics and input/output statistics for devices, partitions and network filesystems.
- [mpstat\(1\)](#) reports individual or combined processor related statistics.
- [pidstat\(1\)](#) reports statistics for Linux tasks (processes) : I/O, CPU, memory, etc.
- [sar\(1\)](#) collects, reports and saves system activity information (CPU, memory, disks, interrupts, network interfaces, TTY, kernel tables, etc.)
- [sadc\(8\)](#) is the system activity data collector, used as a backend for sar.
- [sa1\(8\)](#) collects and stores binary data in the system activity daily data file. It is a front end to sadc designed to be run from cron.
- [sa2\(8\)](#) writes a summarized daily activity report. It is a front end to sar designed to be run from cron.
- [sadf\(1\)](#) displays data collected by sar in multiple formats (CSV, XML, etc.) This is useful to load performance data into a database, or import them in a spreadsheet to make graphs.
- [sysstat\(5\)](#) is just a manual page for sysstat configuration file, giving the meaning of environment variables used by sysstat commands.
- [nfsiostat-sysstat\(1\)](#) reports input/output statistics for network filesystems (NFS).
- [cifsioatstat\(1\)](#) reports CIFS statistics.

Go to the [Features](#) page to display a list of sysstat's main features, and to the [Matrix of activities](#) to list all the possible activities for sar and corresponding options to use with sar and sadc.”

<http://sysadmin.lk/howto-install-sar-ksar-system-activity-info-centos-redhat/>

[sar Tutorial](#)

kSar: GUI front-end

<https://grepora.com/2018/01/17/ksar-generating-graphs-from-sar-reports/>

[Collect and report Linux System Activity Information with sar](#)

Fedora-specific:

<https://computersecuritystudent.com/UNIX/FEDORA/lesson13/index.html>

Useful:

[“10 Useful Sar \(Sysstat\) Examples for UNIX / Linux Performance Monitoring
by RAMESH NATARAJAN “](#)

Useful:

See the [Matrix of activities](#) – essentially, a table of which activities sar can keep track of (with the option switch mentioned).

One can always install the distributor version of sysstat, using the appropriate package manager. On Debian/Ubuntu/Mint:

```
sudo apt-get install sysstat
```

Below, we download and install by hand the latest stable version of sysstat.

[Download](#)

As of 13 February 2015, the latest stable version of sar is 11.0.3.

From: [“10 Useful Sar \(Sysstat\) Examples for UNIX / Linux Performance Monitoring
by RAMESH NATARAJAN “](#)

When manually configuring sar, make sure to use:

```
./configure --enable-install-cron
```

Note: Make sure to pass the option `--enable-install-cron`. This does the following automatically for you. If you don't configure sysstat with this option, you have to do this ugly job yourself manually.

```
# cat /etc/issue
Ubuntu 14.10 \n \l

# uname -a
Linux kaiwan-ThinkPad-X220 3.16.0-31-generic #43-Ubuntu SMP Tue Mar 10
17:37:36 UTC 2015 x86_64 x86_64 x86_64 GNU/Linux
# ls -l /usr/local/lib/sa/          << The sa[12] & sadc is
                                   installed under
/usr/local/lib/sa >>
total 92
-rwxr-xr-x 1 root root  980 Mar 13 18:26 sa1*
-rwxr-xr-x 1 root root 1391 Mar 13 18:26 sa2*
-rwxr-xr-x 1 root root 84040 Mar 13 18:26 sadc*
#
```

Collect the sar statistics using cron job – sa1 and sa2

Create sysstat file under `/etc/cron.d` directory that will collect the historical sar data.

```
# vi /etc/cron.d/sysstat
*/10 * * * * root /usr/local/lib/sa/sa1 1 1
53 23 * * * root /usr/local/lib/sa/sa2 -A
#
```

If you've installed sysstat from source, the default location of sa1 and sa2 is `/usr/local/lib/sa`. If you've installed using your distribution update method (for example: yum, up2date, or apt-get), this might be `/usr/lib/sa/sa1` and `/usr/lib/sa/sa2`.

<< Older output below >>

(On my Ubuntu 12.04 LTS box) sa1 and sa2 files are here:
`/usr/lib/sysstat`

Changes `ENABLED="true"` in `/etc/default/sysstat`.

Cron entries:
`/etc/cron.d/sysstat`

Once the cron entry is enabled and running, use 'sar' to see stuff about system state.
 It uses the (daily) logs generated here:

```
# ls -l /var/log/sysstat/
total 224K
-rw-r--r-- 1 root root 177K Aug  2 19:10 sa02
-rw-r--r-- 1 root root  36K Aug  3 10:22 sa03
#
```

(See url above for lots of examples)

Eg.

To see disk IO:

```
# sar -b
Linux 3.2.21 (kaiwan-ThinkPad-X220)   Thursday 02 August 2012   _i686_
(4 CPU)

05:47:01  IST          tps          rtps          wtps        bread/s        bwrtn/s
05:48:01  IST          6.71          3.93          2.78         33.32         385.74
05:50:01  IST          3.48          1.05          2.43         31.46         47.38
05:55:01  IST          6.61          2.86          3.75        107.83         70.16
...
06:15:01  IST          4.03          1.42          2.62        318.56        125.16
06:16:01  IST          3.32          0.23          3.08          3.33        399.67
06:17:01  IST          2.68          0.05          2.63          0.80         49.58
Average:           6.05          3.01          3.04        104.66        102.56
#
```

<< *the time interval to capture stats was initially set to 5 min; later made it 1 min* >>

Eg to see memory used/freed/cached etc

```
# sar -r
Linux 3.2.21 (kaiwan-ThinkPad-X220)   Thursday 02 August 2012   _i686_   (4 CPU)

05:47:01  IST kmemfree kmemused %memused kbbuffers kbcached kbcommit %commit kbactive kbinact
05:48:01  IST   215756  2757640   92.74    66376    710088   6015044    67.04   1800532   748684
05:50:01  IST   219772  2753624   92.61    66632    710552   5988884    66.75   1796152   749388
05:55:01  IST   189812  2783584   93.62    67432    738436   6003928    66.92   1808652   767520
...
06:17:01  IST   398764  2574632   86.59    53092    649132   5801668    64.66   1690800   685692
06:18:01  IST   401388  2572008   86.50    53252    645672   5795812    64.60   1690968   682312
06:19:01  IST   396728  2576668   86.66    53400    650764   5800700    64.65   1690992   687532
Average:   343353  2630043   88.45    53914    661357   5904121    65.81   1729417   700077
#
```

Eg. Use “sar -B” to generate paging statistics. i.e Number of KB paged in (and out) from disk per second.

```
# sar -B
Linux 3.2.21 (kaiwan-ThinkPad-X220)   Thursday 02 August 2012   _i686_   (4 CPU)

05:47:01  IST   pgpgin/s pgpgout/s   fault/s   majflt/s   pgfree/s pgscank/s pgscand/s pgsteal/s   %vmeff
05:48:01  IST    16.66   192.87   154.03     0.07    383.17     0.00     0.00     0.00     0.00
05:50:01  IST    15.73    23.69   186.56     0.23    299.08     0.00     0.00     0.00     0.00
05:55:01  IST    53.92    35.08  1057.40     0.48   1205.52     0.00     0.00     0.00     0.00
...
06:19:01  IST     3.13    27.46   545.74     0.02    692.15     0.00     0.00     0.00     0.00
06:20:01  IST     1.33    25.72   276.02     0.02    330.16     0.00     0.00     0.00     0.00
06:21:01  IST     0.33    25.93   547.96     0.02    672.74     0.00     0.00     0.00     0.00
Average:   46.32   48.25   608.41     0.45   787.93    27.55     0.00    20.69    75.11
#
```

etc etc!

iostat

Display :

- c Display the CPU utilization report.
- d n Display the device utilization report., every n seconds
- x Display extended statistics.
- z Tell iostat to omit output for any devices for which there was no activity during the sample period.

```
$ iostat -c -d 3 -x -z
```

```
Linux 3.13.0-24-generic (kaiwan-desktop) Wednesday 24 December 2014
_x86_64_ (4 CPU)
```

```
avg-cpu:  %user   %nice %system %iowait  %steal   %idle
           17.53    0.00    4.09    5.84    0.00   72.54
```

<< From 'man iostat':
CPU Utilization Report

The first report generated by the iostat command is the CPU Utilization Report. For multiprocessor systems, the CPU values are global averages among all processors. The report has the following format:

Show the percentage of

%user : % CPU utilization that occurred while executing at the user level (application).
%nice : % CPU utilization that occurred while executing at the user level with nice priority.
%system : % CPU utilization that occurred while executing at the system level (kernel).
%iowait : % time that the CPU or CPUs were idle during which the system had an outstanding disk I/O request.
%steal : % time spent in involuntary wait by the virtual CPU or CPUs while the hypervisor was servicing another virtual processor.
%idle : % time that the CPU or CPUs were idle and the system did not have an outstanding disk I/O request.
 >>

```
Device: rrqm/s wrqm/s r/s w/s rkB/s  wkB/s  avgrq-sz avgqu-sz   await  r_await
w_await svctm  %util
sda      0.00    0.00 115.33  7.67 1100.00 229.33   21.62    0.45    3.69    3.78   2.26
1.99    4.53
```

<< From 'man iostat':

Device: This column gives the device (or partition) name as listed in the /dev directory.

rrqm/s : The number of read requests merged per second that were queued to the device.
 wrqm/s : The number of write requests merged per second that were queued to the device.
 r/s : The number (after merges) of read requests completed per second for the device.
 w/s : The number (after merges) of write requests completed per second for the device.
 rsec/s (rkB/s, rMB/s) : The number of sectors (kilobytes, megabytes) read from the device per second.
 wsec/s (wkB/s, wMB/s) : The number of sectors (kilobytes, megabytes) written to the device per second.
 avgrq-sz : The average size (in sectors) of the requests that were issued to the device.
 avgqu-sz : The average queue length of the requests that were issued to the device.
 await : The average time (in milliseconds) for I/O requests issued to the device to be served. This includes the time spent by the requests in queue and the time spent servicing them.
 r_await : The average time (in milliseconds) for read requests issued to the device to be served. This includes the time spent by the requests in queue and the time spent servicing them.
 w_await : The average time (in milliseconds) for write requests issued to the device to be served. This includes the time spent by the requests in queue and the time spent servicing them.
 svctm : The average service time (in milliseconds) for I/O requests that were issued to the device. **Warning!**
 Do not trust this field any more. This field will be removed in a future sysstat version.

%util : Percentage of CPU time during which I/O requests were issued to the device (bandwidth utilization for the device). Device saturation occurs when this value is close to 100%.

>>

...

```
avg-cpu:  %user   %nice %system %iowait  %steal   %idle
           25.17    0.00    1.83    0.58    0.00   72.42
```

```
Device:            rrqm/s   wrqm/s     r/s     w/s    rkB/s    kB/s avgrq-sz avgqu-sz
await r_await w_await  svctm  %util
sda              0.00    8.67    4.00    6.00    73.33  1449.33   304.53    0.02
2.13    0.00    3.56    2.00    2.00
```

```
avg-cpu:  %user   %nice %system %iowait  %steal   %idle
           24.96    0.00    2.25    0.83    0.00   71.95
```

```
Device:            rrqm/s   wrqm/s     r/s     w/s    rkB/s    kB/s avgrq-sz avgqu-sz
await r_await w_await  svctm  %util
sda              0.00    5.00    4.67    5.67   110.67  1470.67   306.06    0.02
1.68    0.29    2.82    1.29    1.33
```

```
avg-cpu:  %user   %nice %system %iowait  %steal   %idle
           25.88    0.00    2.34    0.83    0.00   70.95
```

```
Device:            rrqm/s   wrqm/s     r/s     w/s    rkB/s    kB/s avgrq-sz avgqu-sz
await r_await w_await  svctm  %util
sda              0.00    5.00    4.67   18.33    41.33  1721.33   153.28    0.01
0.64    0.00    0.80    0.46    1.07
```

```
avg-cpu:  %user   %nice %system %iowait  %steal   %idle
           8.02    0.00    1.25    0.84    0.00   89.89
```

```
Device:            rrqm/s   wrqm/s     r/s     w/s    rkB/s    kB/s avgrq-sz avgqu-sz
await r_await w_await  svctm  %util
sda              0.00    7.00    1.67    3.00    13.33   401.33   177.71    0.03
5.43    3.20    6.67    5.43    2.53
```

...

#

Related: pidstat(1), mpstat(1), vmstat(8)

blktrace

See the man page, with examples.

Interpreting Raw Statistics from the Kernel

The Linux kernel stores disk accounting information and stats within both the /proc and /sys pseudo-file systems.

Documentation: <https://www.kernel.org/doc/Documentation/iostats.txt>

I/O statistics fields

Since 2.4.20 (and some versions before, with patches), and 2.5.45, more extensive disk statistics have been introduced to help measure disk activity. Tools such as sar and iostat typically interpret these and do the work for you, but in case you are interested in creating your own tools, the fields are explained here.

In 2.4 now, the information is found as additional fields in /proc/partitions. In 2.6, the same information is found in two places: one is in the file /proc/diskstats, and the other is within the sysfs file system, which must be mounted in order to obtain the information. Throughout this document we'll assume that sysfs is mounted on /sys, although of course it may be mounted anywhere. Both /proc/diskstats and sysfs use the same source for the information and so should not differ.

Here are examples of these different formats:

2.4:

```

 3      0    39082680 hda 446216 784926 9550688 4382310 424847 312726 5922052
19310380 0 3376340 23705160
 3      1     9221278 hda1 35486 0 35496 38030 0 0 0 0 0 38030 38030
```

2.6 sysfs:

```

446216 784926 9550688 4382310 424847 312726 5922052 19310380 0 3376340
23705160
35486      38030      38030      38030
```

2.6 diskstats:

```

 3      0    hda 446216 784926 9550688 4382310 424847 312726 5922052 19310380 0
3376340 23705160
 3      1    hda1 35486 38030 38030 38030
```

On 2.4 you might execute "grep 'hda ' /proc/partitions". On 2.6, you have a choice of "cat /sys/block/hda/stat" or "grep 'hda ' /proc/diskstats". The advantage of one over the other is that the sysfs choice works well if you are watching a known, small set of disks. /proc/diskstats may be a better choice if you are watching a large number of disks because you'll avoid the overhead of 50, 100, or 500 or more opens/closes with each snapshot of your disk statistics.

In 2.4, the statistics fields are those after the device name. In the above example, the first field of statistics would be 446216.

By contrast, in 2.6 if you look at `/sys/block/hda/stat`, you'll find just the eleven fields, beginning with 446216. If you look at `/proc/diskstats`, the eleven fields will be preceded by the major and minor device numbers, and device name.

Each of these formats provides eleven fields of statistics, each meaning exactly the same things. All fields except field 9 are cumulative since boot. Field 9 should go to zero as I/Os complete; all others only increase (unless they overflow and wrap). Yes, these are (32-bit or 64-bit) unsigned long (native word size) numbers, and on a very busy or long-lived system they may wrap. Applications should be prepared to deal with that; unless your observations are measured in large numbers of minutes or hours, they should not wrap twice before you notice them.

[P.T.O.]

Each set of stats only applies to the indicated device; if you want system-wide stats you'll have to find all the devices and sum them all up.

Field 1 -- # of reads completed

This is the total number of reads completed successfully.

Field 2 -- # of reads merged, field 6 -- # of writes merged

Reads and writes which are adjacent to each other may be merged for efficiency. Thus two 4K reads may become one 8K read before it is ultimately handed to the disk, and so it will be counted (and queued) as only one I/O. This field lets you know how often this was done.

Field 3 -- # of sectors read

This is the total number of sectors read successfully.

Field 4 -- # of milliseconds spent reading

This is the total number of milliseconds spent by all reads (as measured from `__make_request()` to `end_that_request_last()`).

**Reads : fields 1-4.
Cumulative**

**Writes : fields 5-8.
Cumulative**

Field 5 -- # of writes completed

This is the total number of writes completed successfully.

Field 6 -- # of writes merged

See the description of field 2.

Field 7 -- # of sectors written

This is the total number of sectors written successfully.

Field 8 -- # of milliseconds spent writing

This is the total number of milliseconds spent by all writes (as measured from `__make_request()` to `end_that_request_last()`).

Field 9 -- # of I/Os currently in progress

The only field that should go to zero. Incremented as requests are given to appropriate struct request_queue and decremented as they finish.

Field 10 -- # of milliseconds spent doing I/Os

This field increases so long as field 9 is nonzero.

Field 11 -- weighted # of milliseconds spent doing I/Os

This field is incremented at each I/O start, I/O completion, I/O merge, or read of these stats by the number of I/Os in progress (field 9) times the number of milliseconds spent doing I/O since the last update of this field. This can provide an easy measure of both I/O completion time and the backlog that may be accumulating.

	Reads completed	Reads merged	Sectors read	# ms reading				# I/Os in progress	# ms doing I/Os	Weighted #ms doing I/Os	
<< Eg. # while [true] ; do	1	2	3	4	5	6	7	8	9	10	11
Field #	1	2	3	4	5	6	7	8	9	10	11
	913078	20345	27130450	18524332	542206	382325	17047792	70554076	0	9383032	89078332
	913091	20345	27130594	18524444	542206	382325	17047792	70554076	0	9383144	89078444
	913091	20345	27130594	18524444	542206	382325	17047792	70554076	0	9383144	89078444
	913091	20345	27130594	18524444	542206	382325	17047792	70554076	0	9383144	89078444
	913099	20345	27130698	18524580	542206	382325	17047792	70554076	0	9383188	89078580
								</			

```
--snip--
913657 20345 27140066 18537292 542313 382483 17058968 70555996 0 9385604 89093208
913661 20345 27140322 18537312 542313 382483 17058968 70555996 30 9385632 89093468
913783 20345 27184866 18539980 542386 382485 17094968 70558292 2 9386636 89098200
913888 20345 27237610 18541864 542445 382485 17152312 70561196 2 9387636 89103068
914000 20345 27294442 18543664 542501 382485 17206328 70562856 6 9388640 89107144
914114 20345 27351794 18545820 542558 382485 17260856 70564716 8 9389640 89110568
914237 20347 27411954 18548380 542626 382515 17324600 70567356 2 9390828 89115668
914316 20347 27447530 18550152 542682 382522 17360648 70570596 0 9391688 89120664
914316 20347 27447530 18550152 542682 382522 17360648 70570596 0 9391688 89120664
^C
#
>>
```

To avoid introducing performance bottlenecks, no locks are held while modifying these counters. This implies that minor inaccuracies may be introduced when changes collide, so (for instance) adding up all the read I/Os issued per partition should equal those made to the disks ... but due to the lack of locking it may only be very close.

```
--snip--
```

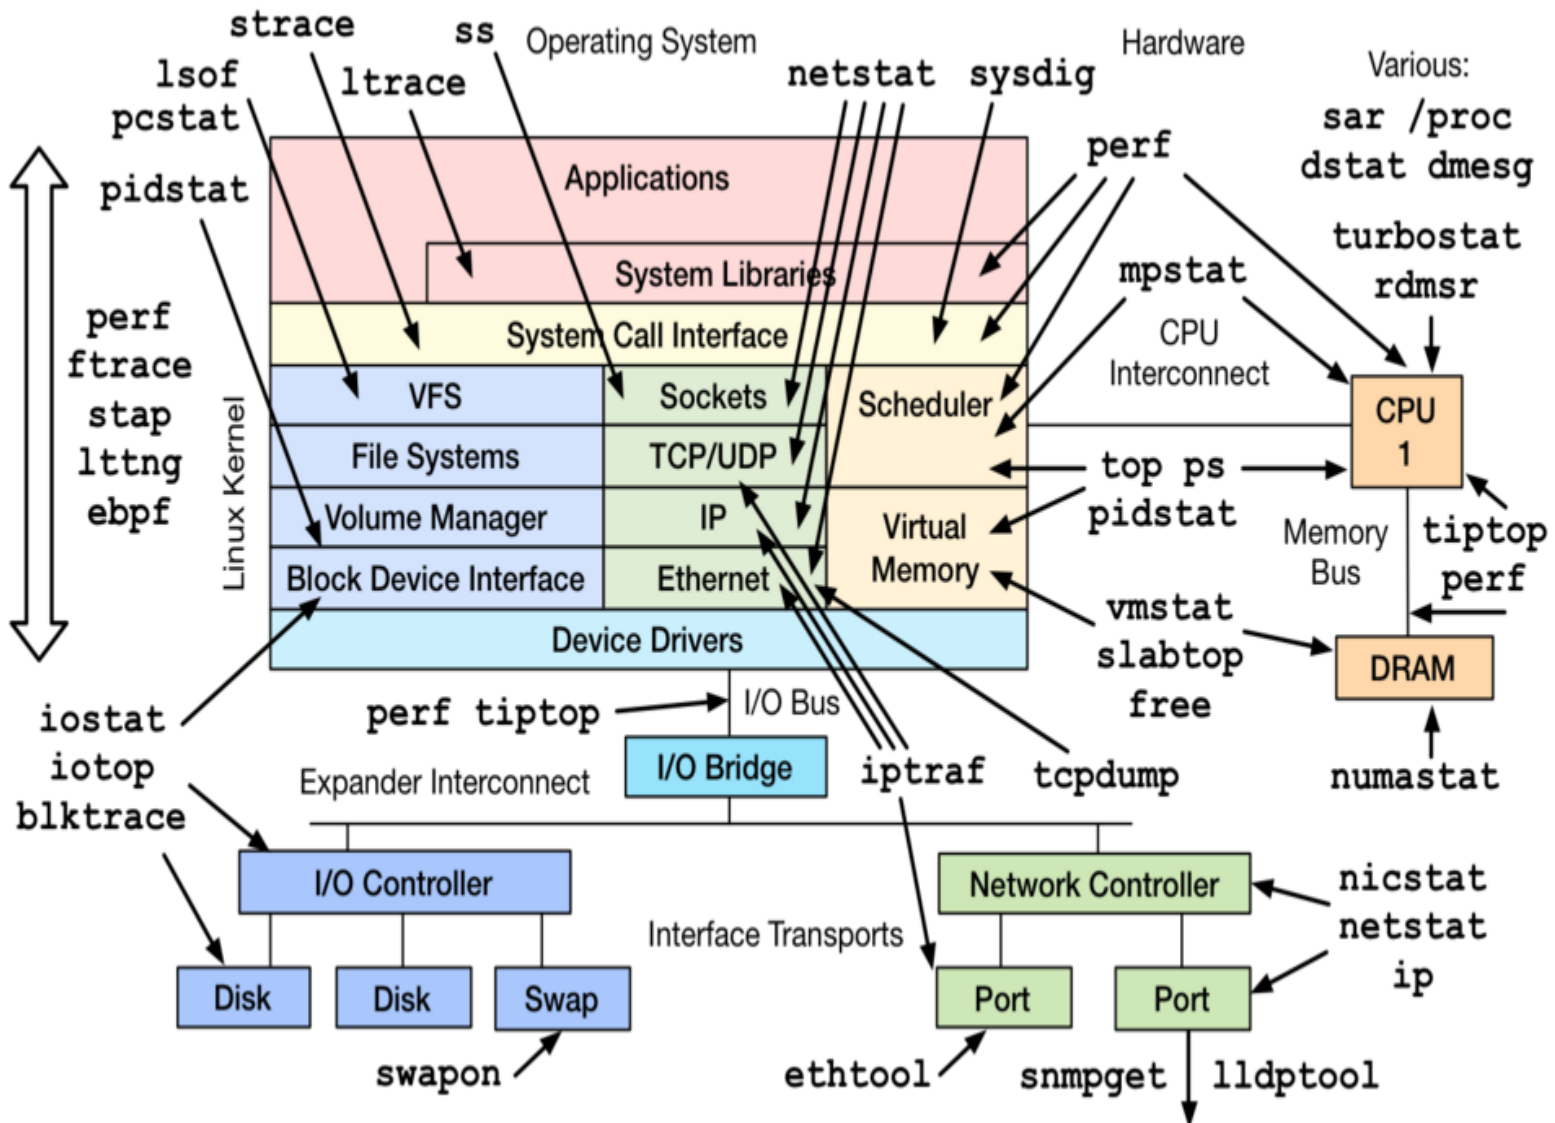
In 2.6, all disk statistics were removed from /proc/stat. In 2.4, they appear in both /proc/partitions and /proc/stat, although the ones in /proc/stat take a very different format from those in /proc/partitions (see proc(5), if your system has it.)

```
-- ricklind@us.ibm.com
```

[Source](#)

Linux Performance Observability Tools, Brendan Gregg

Linux Performance Observability Tools

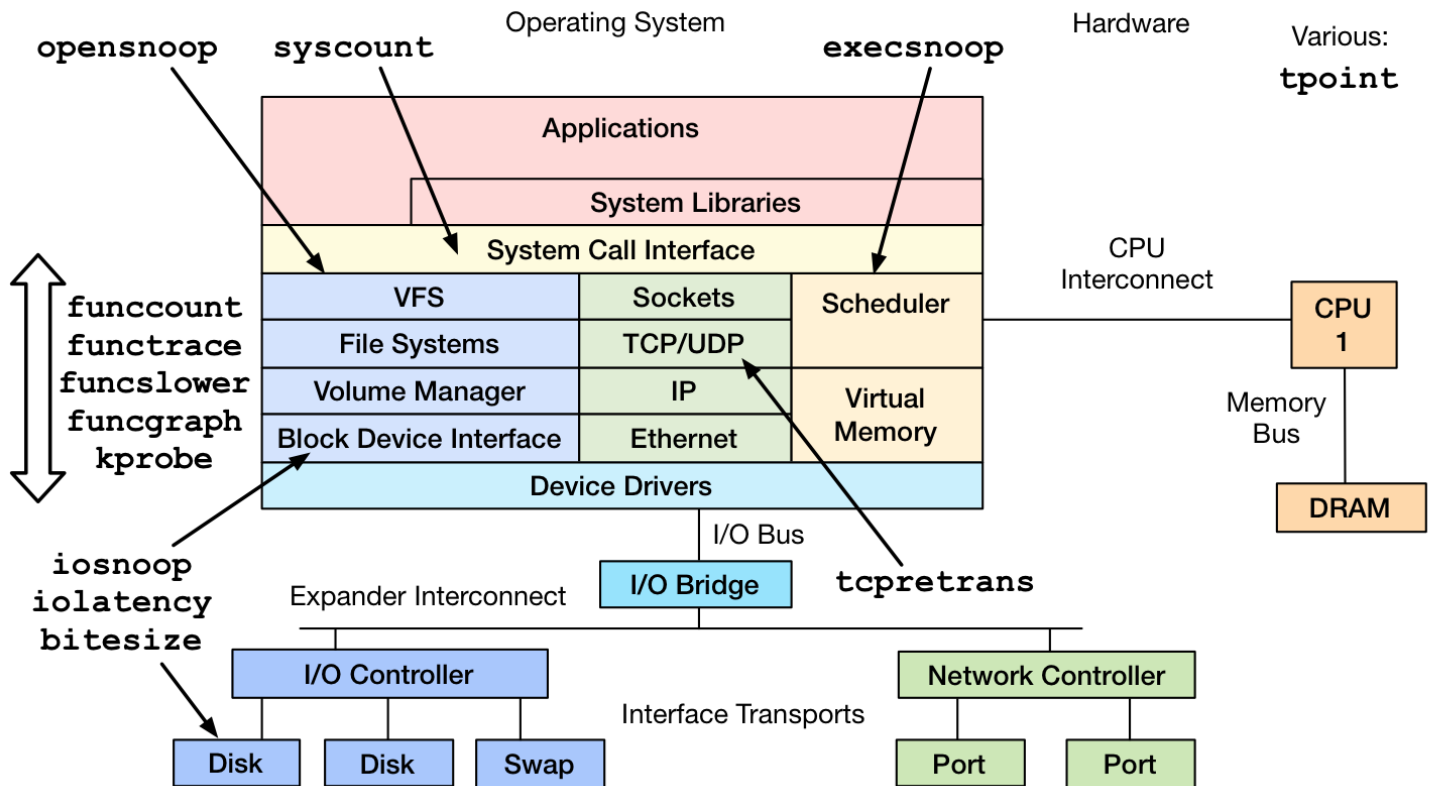


<http://www.brendangregg.com/linuxperf.html> 2015

[Source](#)

Linux Performance Observability Tools: Perf Tools, Brendan Gregg

Linux Performance Observability Tools: perf-tools



<https://github.com/brendangregg/perf-tools#contents>

!WARNING! Many of these tools are considered to be Experimental and Unstable!

<< **Optional / FYI** >>

Network performance

Measure the quality and throughput of a network link using an excellent, open-source, scriptable & easy-to-use tool iPerf. [A good Iperf tutorial is available here.](#)

Source: Essential Linux Device Drivers, S Venkateshwaran

--snip--

Network Throughput

Several tools are available to benchmark network performance. Netperf, available for free from www.netperf.org, can set up complex TCP/UDP connection scenarios. You can use scripts to control characteristics such as protocol parameters, number of simultaneous sessions, and size of data blocks.

Benchmarking is accomplished by comparing the resulting throughput with the maximum practical bandwidth that the networking technology yields. For example, a 155Mbps ATM adapter produces a maximum IP throughput of 135Mbps, taking into account the ATM cell header size, overheads due to the ATM Adaptation Layer (AAL), and the occasional maintenance cells sent by the physical Synchronous Optical Networking (SONET) layer.

To obtain optimal throughput, you have to design your NIC driver for high performance. In addition, you need an in-depth understanding of the network protocol that your driver ferries.

...

FYI: Commands to spew out hardware-related information:

- procfs
 - sysfs
 - lshw
 - lscpu
 - lspci
 - lsusb
 - dmidecode [x86]
-

[Source](#)

CPU Workload Characterization


Q. What's causing CPU load?

- Who?
- How?
- Why?
- What?

A.

- Who? : top, htop
 - How? : system monitoring tools
 - Why? : perf record -g flame graphs
 - What? : perf stat -a -d
-




Linux Operating System Specialized



The highest quality Training on:

Linux Fundamentals, CLI and Scripting
Linux Systems Programming
Linux Kernel Internals
Linux Device Drivers
Embedded Linux
Linux Debugging Techniques
New! Linux OS for Technical Managers

Please do visit our website for details:
<http://kaiwantech.in>



<http://kaiwantech.in>

kaiwanTECH Linux OS Corporate Training Programs

Please do check out our current offering of world-class, seriously-valuable, high on returns, technical Linux OS corporate training programs here: <http://bit.ly/ktcorp>